


NumLin: Linear Types for Linear Algebra

Dhruv C. Makwana 

Unaffiliated dhruvmakwana.com

dcm41@cam.ac.uk

Neelakantan R. Krishnaswami 

Department of Computer Science and Technology, University of Cambridge, United Kingdom

nk480@cl.cam.ac.uk

Abstract

We present NUMLIN, a functional programming language designed to express the APIs of low-level linear algebra libraries (such as BLAS/LAPACK) safely and explicitly, through a brief description of its key features and several illustrative examples. We show that NUMLIN's type system is sound and that its implementation improves upon naïve implementations of linear algebra programs, almost towards C-levels of performance. Lastly, we contrast it to other recent developments in linear types and show that using linear types and fractional permissions to express the APIs of low-level linear algebra libraries is a simple and effective idea.

2012 ACM Subject Classification Theory of computation → Program specifications

Keywords and phrases numerical, linear, algebra, types, permissions, OCaml

Digital Object Identifier 10.4230/LIPIcs.CVIT.2016.23

Supplement Material www.github.com/dc-mak/lt4la

1 Introduction

Programmers writing numerical software often find themselves caught on the horns of a dilemma. The foundational, low-level linear algebra libraries such as BLAS and LAPACK offer programmers very precise control over the memory lifetime and usage of vector and matrix values. However, this power comes paired with the responsibility to manually manage the memory associated with each array object, and in addition to bringing in the familiar difficulties of reasoning about lifetimes, aliasing and sharing that plague low-level systems programming, this also moves the APIs away from the linear-algebraic, mathematical style of thinking that numerical programmers want to use.

As a result, programmers often turn to higher-level languages such as Matlab, R and Numpy, which offer very high level array abstractions that can be viewed as ordinary mathematical values. This makes programming safer, as well as making prototyping and verification much easier, since it lets programmers write programs which bear a closer resemblance to the formulas that the mathematicians and statisticians designing these algorithms prefer to work with, and ensures that program bugs will reflect incorrectly-computed values rather than heap corruption.

The intention is that these languages can use libraries BLAS and LAPACK, without having to expose programmers to explicit memory management. However, this benefit comes at a price: because user programs do not worry about aliasing, the language implementations cannot in general exploit the underlying features of the low-level libraries that let them explicitly manage and reuse memory. As a result, programs written in high-level statistical languages can be much less memory-efficient than programs that make full use of the powers the low-level APIs offer.

So in practice, programmers face a tradeoff: they can eschew safety and exploit the full power of the underlying linear algebra libraries, or they can obtain safety at the price of



© Dhruv C. Makwana and Neelakantan R. Krishnaswami;
licensed under Creative Commons License CC-BY

42nd Conference on Very Important Topics (CVIT 2016).

Editors: John Q. Open and Joan R. Access; Article No. 23; pp. 23:1–23:43

Leibniz International Proceedings in Informatics



LIPICs Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany

unnneeded copies and worse memory efficiency. In this work, show that this tradeoff is not a fundamental one.

NUMLIN is a functional programming language designed to express the APIs of low-level linear algebra libraries (such as BLAS/LAPACK) safely and explicitly. It does so by combining linear types, fractional permissions, runtime errors and recursion into a small, easily understandable, yet expressive set of core constructs.

NUMLIN allows a novice to understand and work with complicated linear algebra library APIs, as well as point out subtle aliasing bugs and reduce memory usage in existing programs. In fact, we were able to use NUMLIN to find linearity and aliasing bugs in a linear algebra algorithm that was *generated* by another program *specifically designed to translate matrix expressions into an efficient sequence of calls to linear algebra routines*. We were also able to reduce the number of temporaries used by the same algorithm, using NUMLIN’s type system to guide us.

NUMLIN’s implementation supports several syntactic conveniences as well as a *usable* integration with real OCaml libraries.

1.1 Contributions

In this paper

- we describe NUMLIN, a linearly typed language for linear algebra programs
- we illustrate that NUMLIN’s design and features are well-suited to its intended domain with progressively sophisticated examples
- we prove NUMLIN’s soundness, using a step-indexed logical relation
- we describe a very simple, unification based type-inference algorithm for polymorphic fractional permissions (similar to ones used for parametric polymorphism), demonstrating an alternative approach to dataflow analysis [5]
- we describe an implementation that is both compatible with and usable from existing code
- we show an example of how using NUMLIN helped highlight linearity and aliasing bugs, and reduce the memory usage of a *generated* linear algebra program
- we show that using NUMLIN, we can achieve parity with C for linear algebra routines, whilst having much better static guarantees about the linearity and aliasing behaviour of our programs.

2 NumLin Overview and Examples

2.1 Overview

The core type theory of NUMLIN is a nearly off-the-shelf linear type theory, supporting familiar features such as linear function spaces $A \multimap B$ and tensor products $A \otimes B$. We adopt linearity – the restriction that each program variable be used at most once – since it allows us to express purely functional APIs for numerical library routines that mutate arrays and matrices [17]. Due to linearity, values cannot alias and are only used once, which means that linearly-typed updates result in no *observable* mutation.

As a result, programmers can reason about NUMLIN expressions as if they were ordinary mathematical expressions – as indeed they are! We are merely adopting a stricter type discipline than usual to make managing memory safe.

2.1.1 Intuitionism: ! and Many

However, linearity by itself is not sufficient to produce an expressive enough programming language. For values such as booleans, integers, floating-point numbers as well as pure functions, we need to be able to use them *intuitionistically*, that is, more than once or not at all. For this reason, we have the ! constructor at the type level and its corresponding Many constructor and let Many <id> = .. in .. eliminator at the term level. Because we want to restrict how a programmer can alias pointers and prevent a programmer from ignoring them (a memory leak), NUMLIN enforces simple syntactic restrictions on which values can be wrapped up in a Many constructor (details in Section 3).

2.1.2 Fractional Permissions

There are also valid cases in which we would want to alias pointers to a matrix. The most common is exemplified by the BLAS routine `gemm`, which (rather tersely) stands for *GEneric Matrix Multiplication*. A simplified definition of `gemm`(α , A, B, β , C) is $C := \alpha AB + \beta C$. In this case, A and B may alias each other but neither may alias C, because it is being written to. Related to *mutating* arrays and matrices is *freeing* them. Here, we would also wish to restrict aliasing so that we do not free one alias and then attempt to use another. Although linearity on its own suffices to prevent use-after-free errors when values are *not* aliased (a freed value is *out of scope* for the rest of the expression), we still need another simple, yet powerful concept to provide us with the extra expressivity of aliasing *without* losing any of the benefits of linearity.

Fractional permissions provide exactly this. Concretely, types of (pointers to) arrays and matrices are *parameterised* by a *fraction*. A fraction is either 1 (2^0) or exactly *half* of another fraction (2^{-k} , for natural k). The former represents complete ownership of that value: the programmer may mutate or free that value as they choose; the latter represents read-only access or a *borrow*: the programmer may read from the value but not write to or free it. Creating an array/matrix gives you ownership of it, so too does having one (with a fractional permission of 2^0) passed in as an argument.

In NUMLIN, we can produce two aliases of a single array/matrix, by *sharing* it. If the original alias had a fractional permission of 2^{-k} then the two new aliases of it will have a fractional permission of $2^{-(k+1)}$ each. Thanks to linearity, the original array/matrix with a fractional permission of 2^{-k} will be out of scope after the sharing. When an array/matrix is shared as such, we can prevent the programmer from freeing or mutating it by making the types of `free` and `set` (for mutation) require a *whole* (2^0) permission.

If we have two aliases *to the same matrix* with *identical* fractional permissions ($2^{-(k+1)}$), we can recombine or *unshare* them back into a single one, with a larger 2^{-k} permission. As before, thanks to linearity, the original two aliases will be out of scope after unsharing.

2.1.3 Runtime Errors

Aside from out-of-bounds indexing, matrix unsharing is one of only *two* operations that can fail at runtime (the other being dimension checks, such as for `gemm`). The check being performed is a simple sanity check that the two aliasing pointers passed to `unshare` point to the same array/matrix. Section 5 contains an overview of how we could remove the need for this by tracking pointer identities statically by augmenting the type system further.

```

129 let rec factorial ( !x : !int ) : !int =
130   if x < 0 || x = 0 then
131     1
132   else
133     x * factorial (x - 1) in factorial
134 ;;

```

■ **Figure 1** Factorial function in NUMLIN.

129 2.1.4 Recursion

130 The final feature of NUMLIN which makes it sufficiently expressive is recursion (and of
 131 course, conditional branches to ensure termination). Conditional branches are implemented
 132 by ensuring that both branches use the same set of linear values. A function can be recursive
 133 if it captures no linear values from its environment. Like with **Many**, this is enforced via
 134 simple syntactic restrictions on the definition of recursive functions.

135 2.2 Examples

136 2.2.1 Factorial

137 Although a factorial function (Figure 1) may seem like an aggressively pedestrian first example,
 138 in a linearly typed language such as NUMLIN it represents the culmination of many features.

139 To simplify the design and implementation of NUMLIN’s type system, recursive functions
 140 must have full type annotations (non-recursive functions need only their argument types
 141 annotated). Its body is a closed expression (with respect to the function’s arguments), so it
 142 type-checks (since it does not capture any linear values from its environment).

143 The only argument is `!x : !int`. The `!` annotation on `x` is a syntactic convenience for
 144 declaring the value to be used intuitionistically, its full and precise meaning is described in
 145 Section 4.1.1.

146 The condition for an `if` may or may not use linear values (here, with `x < 0 || x = 0`, it
 147 does not). Any linear values used by the condition would not be in scope in either branch of
 148 the `if`-expression. Both branches use `x` differently: one ignores it completely and the other
 149 uses it twice.

150 All numeric and boolean literals are implicitly wrapped in a **Many** and all primitives
 151 involving them return a `!int`, `!bool` or `!elt` (types of elements of arrays/matrices, typically
 152 64-bit floating-point numbers). The short-circuiting `||` behaves in exactly the same way as a
 153 boolean-valued `if`-expression.

154 2.2.2 Summing over an Array

155 Now we can add fractional permissions to the mix: Figure 2 shows a simple, tail-recursive
 156 implementation of summing all the elements in an array. There are many new features; first
 157 among them is `!x0 : !elt`, the type of array/matrix elements (64-bit floating point).

158 Second is `('x) (row: 'x arr)` which is an array with a universally-quantified fractional
 159 permission. In particular, this means the body of the function cannot mutate or free the
 160 input array, only read from it. If the programmer did try to mutate or free `row`, then they
 161 would get a helpful error message (Figure 3).

162 Alongside taking a `row: 'x arr`, the function also returns an array with exactly the
 163 same fractional permission as the `row` (which can only be `row`). This is necessary because of
 164 linearity: for the caller, the original array passed in as an argument would be out of scope

```

let rec sum_array (!i : !int) (!n : !int) (!x0 : !elt)
  ('x) (row : 'x arr) : 'x arr * !elt =
  if i = n then
    (row, x0)
  else
    let (row, !x1) = row[i] in
    sum_array (i + 1) n (x0 +. x1) 'x row in
    sum_array
;;

```

■ **Figure 2** Summing over an array in NUMLIN.

```

let row = row[i] := x1 in (* or *) let () = free row in
(* Could not show equality: *)
(* z arr *)
(* with *)
(* 'x arr *)
(* *)
(* Var 'x is universally quantified *)
(* Are you trying to write to/free/unshare an array you don't own? *)
(* In examples/sum_array.lt, at line: 7 and column: 19 *)

```

■ **Figure 3** Attempting to write to or free a read only array in NUMLIN.

165 for the rest of the expression, so it needs to be returned and then rebound to be used for the
 166 rest of the function.

167 An example of this consuming and re-binding is in `let (row, !x1) = row[i]`. Indexing
 168 is implemented as a primitive `get: 'x. 'x arr --o !int --o 'x arr * !elt`. Although
 169 fractional permissions can be passed around explicitly (as done in the recursive call), they
 170 can also be *automatically inferred at call sites*: `row[i] == get _ row i` takes advantage of
 171 this convenience.

172 2.2.3 One-dimensional Convolution

173 Figure 4 extends the set of features demonstrated by the previous examples by mutating one
 174 of the input arrays. A one-dimensional convolution involves two arrays: a read-only kernel
 175 (array of weights) and an input vector. It modifies the input vector *in-place* by replacing
 176 each `write[i]` with a weighted (as per the values in the kernel) sum of it and its neighbours;
 177 intuitively, sliding a dot-product with the kernel across the vector.

178 What's implemented in Figure 4 is a *simplified* version of this idea, so as to not distract
 179 from the features of NUMLIN. The simplifications are:

- 180 ■ the kernel has a length 3, so only the value of `write[i-1]` (prior to modification in the
 181 previous iteration) needs to be carried forward using `x0`
- 182 ■ `write` is assumed to have length `n+1`
- 183 ■ `i`'s initial value is assumed to be 1
- 184 ■ `x0`'s initial value is assumed to be `write[0]`
- 185 ■ the first and last values of `write` are ignored.

186 Mutating an array is implemented similarly to indexing one: a primitive `set: z arr`
 187 `--o !int --o !elt --o z arr`. It consumes the original array and returns a new array
 188 with the updated value. `let written = write[i] := <exp>` is just syntactic sugar for `let`
 189 `written = set write i <exp>`.

```

let rec simp_oned_conv
  (!i : !int) (!n : !int) (!x0 : !elt)
  (write : z arr) ('x) (weights : 'x arr)
  : 'x arr * z arr =
  if n = i then (weights, write) else
  let !w0 <- weights[0] in
  let !w1 <- weights[1] in
  let !w2 <- weights[2] in
  let !x1 <- write[i] in
  let !x2 <- write[i + 1] in
  let written = write[i] := w0 *. x0 +. (w1 *. x1 +. w2 *. x2) in
  simp_oned_conv (i + 1) n x1 written _ weights in
simp_oned_conv
;;

```

■ **Figure 4** *Simplified* one-dimensional convolution.

```

let !square ('x) (x : 'x mat) =
  let (x, (!m, !n)) = sizeM _ x in
  let (x1, x2) = shareM _ x in
  let answer <- new (m, n) [| x1 * x2 |] in
  let x = unshareM _ x1 x2 in
  (x, answer) in
square
;;

```

■ **Figure 5** Linear regression (OLS): $\hat{\beta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$

190 Since `write: z arr` (where `z` stands for $k = 0$, representing a fractional permission
 191 of $2^{-k} = 2^{-0} = 1$), we may mutate it, but since we only need to read from `weights`, its
 192 fractional permission index can be universally-quantified. In the recursive call, we see `_` being
 193 used explicitly to tell the compiler to *infer* the correct fractional permission based on the
 194 given arguments.

195 2.2.4 Squaring a Matrix

196 *The most pertinent aspect of NUMLIN is the types of its primitives.* While the types of
 197 operations such as `get` and `set` might be borderline obvious, the types of BLAS/LAPACK
 198 routines become an *incredibly useful, automated check for using the API correctly.*

199 Figure 5 shows how a linearly-typed matrix squaring function may be written in NUMLIN.
 200 It is a *non-recursive* function declaration (the return type is inferred). Since we would like
 201 to be able to use a function like `square` more than once, it is marked with a `!` annotation
 202 (which also ensures it captures no linear values from the surrounding environment).

203 To square a matrix, first, we extract the dimensions of the argument `x`. Then, because
 204 we need to use `x` twice (so that we can multiply it by itself) but linearity only allows one use,
 205 we use `shareM: 'x. 'x mat --o 'x s mat * 'x s mat` to split the permission `'x` (which
 206 represents 2^{-x}) into two halves (`'x s`, which represents $2^{-(x+1)}$).

207 Even if `x` had type `z mat`, sharing it now enforces the assumption of all BLAS/LAPACK
 208 routines that any matrix which is written to (which, in NUMLIN, is always of type `z mat`)
 209 does not alias any other matrix in scope. So if we did try to use one of the aliases in mutating
 210 way, the expression would not type check, and we would get an error similar to the one in
 211 Figure 3.

```

let !lin_reg ('x) (x : 'x mat)
    ('y) (y : 'y mat) =
  let (x, (!_n, !m)) = sizeM _ x in
  let xy <- new (m, 1) [| xT * y |] in
  let x_T_x <- new (m, m) [| xT * x |] in
  let (to_del, answer) = posv x_T_x xy in
  let () = freeM to_del in
  ((x, y), answer) in
lin_reg
;;

```

■ **Figure 6** Linear regression (OLS): $\hat{\beta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$

212 The line `let answer <- new (m,n) [| x1 * x2 |]` is syntactic sugar for first creating
 213 a new $m \times n$ matrix (`let answer = matrix m n`) and then storing the result of the mul-
 214 tiplication in it (`let ((x1, x2), answer) = gemm 1. _ (x1, false) _ (x2, false) 0.`
 215 `answer`). `false` means the matrix should not be accessed with indices transposed.

216 By using some simple pattern-matching and syntactic sugar, we can:

- 217 ■ write normal-looking, apparently non-linear code
- 218 ■ use matrix expressions directly and have a call to an efficient call to a BLAS/LAPACK
 219 routine inserted with appropriate re-bindings
- 220 ■ retain the safety of linear types with fractional permissions by having the compiler
 221 statically enforce the aliasing and read/write rules implicitly assumed by BLAS/LAPACK
 222 routines.

2.2.5 Linear Regression

223 In Figure 6, we wish to compute $\hat{\beta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$. To do that, first, we extract the
 224 dimensions of matrix \mathbf{x} . Then, we say we would like \mathbf{xy} to be a new matrix, of dimension
 225 $m \times 1$, which contains the result of $\mathbf{X}^T \mathbf{y}$ (using syntactic sugar for `matrix` and `gemm` calls
 226 similar to that used in Figure 5, with a ^T annotation on \mathbf{x} to set \mathbf{x} 's 'transpose indices'-flag
 227 to `true`).

228 However, the line `let x_T_x <- new (m,m) [| xT * x |]`, works for a slightly differ-
 229 ent reason: that pattern is matched to a BLAS call to (`syrk true 1. x 0. x_T_x`), which
 230 only uses \mathbf{x} once. Hence \mathbf{x} can appear *twice* in the *pattern* without any calls to `share`.

231 After computing $\mathbf{x}_T \mathbf{x}$, we need to invert it and then multiply it by \mathbf{xy} . The BLAS
 232 routine `posv: z mat --o z mat --o z mat * z mat` does exactly that: assuming the first
 233 argument is symmetric, `posv` mutates its second argument to contain the desired value. Its
 234 first argument is also mutated to contain the (upper triangular) Cholesky decomposition
 235 factor of the original matrix. Since we do not need that matrix (or its memory) again, we
 236 `free` it. If we forgot to, we would get a `Variable to_del not used` error. Lastly, we return
 237 the `answer` alongside the untouched input matrices (\mathbf{x}, \mathbf{y}) .
 238

2.2.6 L1-Norm Minimisation on Manifolds

239 L1-Norm minimisation is often used in optimisation problems, as a *regularisation* term for
 240 reducing the influence of outliers. Although the below formulation[8] is intended to be used
 241 with *sparse* computations, NUMLIN's current implementation only implements dense ones.
 242 However, it still serves as a useful example of explaining NUMLIN's features.
 243


```

let !l1_norm_min (q : z mat) (u : z mat) =
  let (u, (!_n, !k)) = sizeM _ u in
  let (u, u_T) = transpose _ u in
  let (tmp_n_n, q_inv_u) = gesv q u in
  let i = eye k in
  let to_inv <- [| i + u_T * q_inv_u |] in
  let (tmp_k_k, inv_u_T) = gesv to_inv u_T in
  let () = freeM tmp_k_k in
  let answer <- [| 0. * tmp_n_n + q_inv_u * inv_u_T |] in
  let () = freeM q_inv_u in
  let () = freeM inv_u_T in
  answer in
l1_norm_min
;;

```

■ **Figure 7** L1-norm minimisation on manifolds: $\mathbf{Q}^{-1}\mathbf{U}(\mathbf{I} + \mathbf{U}^T\mathbf{Q}^{-1}\mathbf{U})^{-1}\mathbf{U}^T$

Figure 7 shows even more pattern-matching. Patterns of the form `let <id> <- [| beta * c + alpha * a * b |]` are also desugared to `gemm` calls. Primitives like `transpose`: `'x. 'x mat --o 'x mat * z mat` and `eye`: `!int --o z mat` allocate new matrices; `transpose` returns the transpose of a given matrix and `eye k` evaluates to a $k \times k$ identity matrix.

We also see our first example of re-using memory for different matrices: like with `to_del` and `posv` in the previous example, we do not need the value stored in `tmp_5_5` after the call to `gesv` (a primitive similar to `posv` but for a non-symmetric first argument). However, we can re-use its memory much later to store `answer` with `let answer <- [| 0. * tmp_5_5 + q_inv_u * inv_u_T |]`. Again, thanks to linearity, the identifiers `q` and `tmp_5_5` are out of scope by the time `answer` is bound. Although during execution, all three refer to the same piece of memory, logically they represent different values throughout the computation.

2.2.7 Kalman Filter

A *Kalman Filter*[12] is an algorithm for combining prior knowledge of a state, a statistical model and measurements from (noisy) sensors to produce an estimate a more reliable estimated of the current state. It has various applications (navigation, signal-processing, econometrics) and is relevant here because it is usually presented as a series of complex matrix equations.

Figure 8 shows a NUMLIN implementation of a Kalman filter (equations in Figure 9). A few new features and techniques are used in this implementation:

- `sym` annotations in matrix expressions: when this is used, a call to `symm` (the equivalent of `gemm` but for symmetric matrices so that only half the operations are performed) is inserted
- `copyM_to` is used to re-use memory by *overwriting* the contents of its second argument to that of its first (erroring if dimensions do not match)
- `let new_r <- new [| r_2 |]` creates a copy of `r_2`
- `posvFlip` is like `posv` except for solving $XA = B$
- a lot of memory re-use; the following sets of identifiers alias each other:
 - `r_1`, `r_2` and `k_by_k`
 - `data_1` and `data_2`
 - `mu` and `new_mu`
 - `sigma_hT` and `x`


```

let !kalman
  ('s) (sigma : 's mat) (* n,n *)
  ('h) (h : 'h mat)      (* k,n *)
  (mu : z mat)            (* n,1 *)
  (r_1 : z mat)           (* k,k *)
  (data_1 : z mat)        (* k,1 *) =
  let (h, (!k, !n)) = sizeM _ h in
  (* could use [| sym(sigma) * hT |] but would
     need a (n,k) temporary hT = transpose _ h *)
  let sigma_hT <- new (n, k) [| sigma * h^T |] in
  let r_2 <- [| r_1 + h * sigma_hT |] in
  let (k_by_k, x) = posvFlip r_2 sigma_hT in
  let data_2 <- [| h * mu - data_1 |] in
  let new_mu <- [| mu + x * data_2 |] in
  let x_h <- new (n,n) [| x * h |] in
  let () = freeM (* n,k *) x in
  let sigma2 <- new [| sigma |] in
  let new_sigma <- [| sigma2 - x_h * sym(sigma) |] in
  let () = freeM (* n,n *) x_h in
  ((sigma, h), (new_sigma, (new_mu, (k_by_k, data_2)))) in
kalman
;;

```

■ **Figure 8** Kalman filter: see Figure 9 for the equations this code implements. Line numbers in comments refer to equivalent lines in a C implementation (Figure 17).

$$\mu' = \mu + \Sigma H^T (R + H \Sigma H^T)^{-1} (H \mu - \text{data})$$

$$\Sigma' = \Sigma (I - H^T (R + H \Sigma H^T)^{-1} H \Sigma)$$

■ **Figure 9** Kalman filter equations (credit: matthewrocklin.com).

275 The NUMLIN implementation is much longer than the mathematical equations for two
 276 reasons. First, the NUMLIN implementation is a let-normalised form of the Kalman equations:
 277 since there a large number of unary/binary (and occasionally ternary) sub-expressions in
 278 the equations, naming each one line at a time makes the implementation much longer.
 279 Second, NUMLIN has the additional task of handling explicit allocations, aliasing and frees
 280 of matrices. However, it is exactly this which makes it possible (and often, easy) to spot
 281 additional opportunities for memory re-use. Furthermore, a programmer can explore those
 282 opportunities easily because NUMLIN's type system statically enforces correct memory
 283 management and the aliasing assumptions of BLAS/LAPACK routines.

284 3 Formal System

285 3.1 Core Type Theory

286 The full typing rules are in Appendix A.1, but the key ideas are as follow.

287 A typing judgement consists of $\Theta; \Delta; \Gamma \vdash e : t$.

288 Θ is the environment that tracks which fractional permission variables in scope. Fractional
 289 permissions (the Perm judgement) and types (the Type judgement) are *well-formed* if all of
 290 their free fractional variables are in Θ .

291 Δ is the environment storing non-linearly or *inuitionistically* typed variables.

292 Γ is the environment storing linearly typed variables. Note that rules for typing $()$,

23:10 NumLin: Linear Types for Linear Algebra

booleans, integers and elements are typed with respect to an *empty* linear environment: this means no linear values are needed to produce a value of those types.

$$\frac{}{\Theta; \Delta; \cdot \vdash () : \mathbf{unit}} \quad \text{TY_UNIT_INTRO}$$

Conversely, whenever two or more subexpressions need to be typed, they must consume a disjoint set of linear values (pairs, let-expressions). In the case of if-expressions, both branches must consume the same set of linear values (disjoint to the ones used to evaluate the condition).

$$\frac{\begin{array}{l} \Theta; \Delta; \Gamma \vdash e : \mathbf{!bool} \\ \Theta; \Delta; \Gamma' \vdash e_1 : t' \\ \Theta; \Delta; \Gamma' \vdash e_2 : t' \end{array}}{\Theta; \Delta; \Gamma, \Gamma' \vdash \mathbf{if } e \mathbf{ then } e_1 \mathbf{ else } e_2 : t} \quad \text{TY_BOOL_ELIM}$$

The **Many** introduction and elimination rules are very important. Producing **!**-type values may only be done if the expression inside is a syntactic value which is not a location. This allows all safely duplicable resources, including functions which capture non-linear resources from their environments, but prevents producing aliases of (pointers to) arrays and matrices. This is exactly the same as value-restriction from the world of parametric polymorphism.

$$\frac{\begin{array}{l} \Theta; \Delta; \cdot \vdash v : t \\ v \neq l \end{array}}{\Theta; \Delta; \cdot \vdash \mathbf{Many } v : \mathbf{!}t} \quad \text{TY_BANG_INTRO}$$

Consuming a **!**-type value *moves it* from the linear environment Γ and *into* the intuitionistic environment Δ . This is exactly why **let** $\mathbf{!}x = e_1$ **in** e_2 desugars to **let** **Many** $x = e_1$ **in** **let** **Many** $x = \mathbf{Many } (\mathbf{Many } x)$ **in** e_2 .

$$\frac{\begin{array}{l} \Theta; \Delta; \Gamma \vdash e : \mathbf{!}t \\ \Theta; \Delta, x : t; \Gamma' \vdash e' : t' \end{array}}{\Theta; \Delta; \Gamma, \Gamma' \vdash \mathbf{let } \mathbf{Many } x = e \mathbf{ in } e' : t'} \quad \text{TY_BANG_ELIM}$$

Rules TY_GEN and TY_SPC are for fractional permission generalisation and specialisation respectively. They allow the definition and use of functions that are polymorphic in the fractional permission index of their results and one or more of their arguments.

$$\frac{\Theta, fc; \Delta; \Gamma \vdash e : t}{\Theta; \Delta; \Gamma \vdash \mathbf{fun } fc \rightarrow e : \forall fc. t} \quad \text{TY_GEN} \qquad \frac{\begin{array}{l} \Theta \vdash f \text{ Perm} \\ \Theta; \Delta; \Gamma \vdash e : \forall fc. t \end{array}}{\Theta; \Delta; \Gamma \vdash e[f] : t[f/fc]} \quad \text{TY_SPC}$$

Rule TY_FIX shows how recursive functions are typed. Even though recursive functions are fully annotated, type checking them is interesting for two reasons: to type check the body of the fixpoint, the type of the recursive function is in the *intuitionistic* environment Δ (without this, you would not be able to write a base case) whilst the argument and its type are the *only things in the linear environment* Γ . The latter means that recursive functions can be type checked in an empty environment (thus be wrapped in **Many** and used zero or multiple times).

$$\frac{\Theta; \Delta, g : t \multimap t'; \cdot, x : t \vdash e : t'}{\Theta; \Delta; \cdot \vdash \mathbf{fix}(g, x : t, e : t') : t \multimap t'} \quad \text{TY_FIX}$$

Lastly, types of almost all NUMLIN primitives, as embedded in OCaml's type system, are shown in Appendix G, with some similar ones (like those for binary arithmetic operators) omitted for brevity. The main difference between the OCaml type of a primitive like `gemm` and its NUMLIN counterpart is the inclusion of explicit ' \forall 's. So, `float bang -> ('a mat * bool bang) -> ('b mat * bool bang) -> float bang -> z mat -> ('a mat * 'b mat) * z mat` will correspond to

$$\mathbf{!elt} \multimap \forall x. x \mathbf{mat} \otimes \mathbf{!bool} \multimap \forall y. y \mathbf{mat} \otimes \mathbf{!bool} \multimap \mathbf{!elt} \multimap z \mathbf{mat} \multimap (x \mathbf{mat} \otimes y \mathbf{mat}) \otimes z \mathbf{mat}$$

3.2 Dynamic Semantics

The full, small-step transition relation is in Appendix A.2, but the key ideas are as follow.

Heaps (σ) are multisets containing triples of an abstract location l , a fractional permission f and sized matrices $m_{n,k}$. The notation $l \mapsto_f m_{k_1, k_2}$ should be read as “location l represents f ownership over matrix m (of size $k_1 \times k_2$)”. Each heap-and-expression either steps to another heap-and-expression or a runtime error `err`. In the full grammar definition we see a definition of values and contexts in the language.

We draw the reader's attention to the definitions relating to fractional permissions. Specifically, unlike a lambda, the body of a `fun fc` \rightarrow $_$ must be a syntactic value. The context `fun fc` \rightarrow $[-]$ means expressions can be reduced inside a fractional permission generalisation. This is to emphasize that fractions are merely *compile-time constructs* and do not affect runtime behaviour. Correct usage of fractions is enforced by the type system, so programs do not get stuck. Fractional permissions are specialised using substitution over both the heap and an expression (`OP_FRAC_PERM`).

$$\frac{}{\langle \sigma, (\mathbf{fun} fc \rightarrow v)[f] \rangle \rightarrow \langle \sigma[f c / f], v[f c / f] \rangle} \quad \text{OP_FRAC_PERM}$$

Like with the static semantics, the interesting rules in the dynamic semantics are those relating to primitives. Creating a matrix (`matrix` k_1 k_2) successfully (`OP_MATRIX`) requires non-negative dimensions and returns a (fresh) location of a matrix of those dimensions, extending the heap to reflect that l represents a complete ownership over the new matrix.

$$\frac{0 \leq k_1, k_2 \quad l \text{ fresh}}{\langle \sigma, \mathbf{matrix} k_1 k_2 \rangle \rightarrow \langle \sigma + \{l \mapsto_1 M_{k_1, k_2}\}, l \rangle} \quad \text{OP_MATRIX}$$

Dually, `OP_FREE`, requires a location represent complete ownership before removing it and the matrix it points to from the heap.

$$\frac{}{\langle \sigma + \{l \mapsto_1 m_{k_1, k_2}\}, \mathbf{free} l \rangle \rightarrow \langle \sigma, () \rangle} \quad \text{OP_FREE}$$

Choosing a multiset representation as opposed to a set allows for two convenient invariants: multiplicity of a triple $l \mapsto_f m_{k_1, k_2}$ in the heap corresponds to the number of aliases of l in the expression with type f `mat` and the sum of all the fractions for l will always be 1 (for a closed, well-typed expression). With this in mind, the rules `OP_SHARE` and `OP_UNSHARE_EQ` are fairly natural.

$$\frac{}{\langle \sigma + \{l \mapsto_f m_{k_1, k_2}\}, \mathbf{share}[f] l \rangle \rightarrow \langle \sigma + \{l \mapsto_{\frac{1}{2}f} m_{k_1, k_2}\} + \{l \mapsto_{\frac{1}{2}f} m_{k_1, k_2}\}, (l, l) \rangle} \quad \text{OP_SHARE}$$

$$\frac{}{\langle \sigma + \{l \mapsto_{\frac{1}{2}f} m_{k_1, k_2}\} + \{l \mapsto_{\frac{1}{2}f} m_{k_1, k_2}\}, \mathbf{unshare}[f] \, ll \rangle \rightarrow \langle \sigma + \{l \mapsto_f m_{k_1, k_2}\}, l \rangle} \text{OP_UNSHARE_EQ}$$

Combining all of these features, we see that `OP_GEMM_MATCH` requires that the location being updated (l_3) has complete ownership of over matrix m_3 and can thus change what value it stores to $m_1 m_2 + m_3$. In particular, this places no restriction on l_2 and l_3 : they could be **shared** aliases of the same matrix. Transition rules for other primitives (omitted) follow the same structure: \mapsto_1 for any locations that are written to and \mapsto_{fc} for anything else.

$$\frac{\begin{array}{l} \sigma' \equiv \sigma + \{l_1 \mapsto_{fc_1} m_{1k_1, k_2}\} + \{l_2 \mapsto_{fc_2} m_{2k_2, k_3}\} \\ \sigma_1 \equiv \sigma' + \{l_3 \mapsto_1 m_{3k_1, k_3}\} \\ \sigma_2 \equiv \sigma' + \{l_3 \mapsto_1 (m_1 m_2 + m_3)_{k_1, k_3}\} \end{array}}{\langle \sigma_1, \mathbf{gemm}[fc_1] \, l_1 [fc_2] \, l_2 \, l_3 \rangle \rightarrow \langle \sigma_2, ((l_1, l_2), l_3) \rangle} \text{OP_GEMM_MATCH}$$

3.3 Logical Relation

First, we define an interpretation of heaps with fractional permissions in the style of Bornat et. al [6] (interpreting the multiset as a partial map from locations to the sum of all its associated fractions and a matrix) as well as the n -fold iteration of \rightarrow .

$$\mathcal{H}[\sigma] = \star_{(l, f, m) \in \sigma} [l \mapsto_f m]$$

where

$$(\varsigma_1 \star \varsigma_2)(l) \equiv \begin{cases} \varsigma_1(l) & \text{if } l \in \text{dom}(\varsigma_1) \wedge l \notin \text{dom}(\varsigma_2) \\ \varsigma_2(l) & \text{if } l \in \text{dom}(\varsigma_2) \wedge l \notin \text{dom}(\varsigma_1) \\ (f_1 + f_2, m) & \text{if } (f_1, m) = \varsigma_1(l) \wedge (f_2, m) = \varsigma_2(l) \wedge f_1 + f_2 \leq 1 \\ \text{undefined} & \text{otherwise} \end{cases}$$

We then define a step-indexed logical relation in the style of Morrisett et. al [14]. $(\varsigma, v) \in \mathcal{V}_k[t]$ means it takes a heap with exactly ς resources to produce a value v of type t in at most k steps. So, something like a **unit** or a $!t$ need no resources, whereas a $f \mathbf{mat}$ needs exactly f ownership of a some matrix and a pair needs a \star combination of the heaps required for each component.

$$\begin{aligned} \mathcal{V}_k[\mathbf{unit}] &= \{(\emptyset, *)\} \\ \mathcal{V}_k[f \mathbf{mat}] &= \{(\{l \mapsto_{2-f} _ \}, l)\} \\ \mathcal{V}_k[!t] &= \{(\emptyset, \mathbf{Many} \, v) \mid (\emptyset, v) \in \mathcal{V}_k[t]\} \\ \mathcal{V}_k[t_1 \otimes t_2] &= \{(\varsigma_1 \star \varsigma_2, (v_1, v_2)) \mid (\varsigma_1, v_1) \in \mathcal{V}_k[t_1] \wedge (\varsigma_2, v_2) \in \mathcal{V}_k[t_2]\} \end{aligned}$$

The definition of $\mathcal{V}_k[\forall fc. t]$ says a value and heap must be the same regardless of what fraction is substituted into both; the $k - 1$ is to take into account fraction specialisation takes ones step (`OP_SPC`).

$$\mathcal{V}_k[\forall fc. t] = \{(\varsigma, \mathbf{fun} \, fc \rightarrow v) \mid \forall f. (\varsigma[fc/f], v[fc/f]) \in \mathcal{V}_{k-1}[t[fc/f]]\}$$

To understand the definition of $\mathcal{V}_k[t' \multimap t]$, we must first look at $\mathcal{C}_k[t]$, the computational interpretation of types. Intuitively, it is a combination of a frame rule on heaps (no interference), type-preservation and termination (in $j < k$ steps) to either an error or a

heap-and-expression, with the further condition that if the expression is a syntactic value then it is also one semantically.

$$\mathcal{C}_k[t] = \{(\varsigma_s, e_s) \mid \forall j < k, \sigma_r. \varsigma_s \star \varsigma_r \text{ defined} \Rightarrow \langle \sigma_s + \sigma_r, e_s \rangle \rightarrow^j \mathbf{err} \vee \exists \sigma_f, e_f. \langle \sigma_s + \sigma_r, e_s \rangle \rightarrow^j \langle \sigma_f + \sigma_r, e_f \rangle \wedge (e_f \text{ is a value} \Rightarrow (\varsigma_f \star \varsigma_r, e_f) \in \mathcal{V}_{k-j}[t])\}$$

In this light, $\mathcal{V}_k[t' \multimap t]$ simply says that v is a function and that evaluating the application of it to any argument (of the correct type, requiring its own set of resources, bounded by k steps) satisfies all the aforementioned properties.

$$\mathcal{V}_k[t' \multimap t] = \{(\varsigma_v, v) \mid (v \equiv \mathbf{fun} x : t' \rightarrow e \vee v \equiv \mathbf{fix}(g, x : t', e : t)) \wedge \forall j \leq k, (\varsigma_{v'}, v') \in \mathcal{V}_j[t']. \varsigma_v \star \varsigma_{v'} \text{ defined} \Rightarrow (\varsigma_v \star \varsigma_{v'}, v v') \in \mathcal{C}_j[t]\}$$

The interpretation of typing environments Δ and Γ are with respect to an arbitrary substitution of fractional permissions θ . Note that only the interpretation of Γ involves a (potentially) non-empty heap.

$$\begin{aligned} \mathcal{I}_k[\Delta, x : t]\theta &= \{\delta[x \mapsto v_x] \mid \delta \in \mathcal{I}_k[\Delta]\theta \wedge (\emptyset, v_x) \in \mathcal{V}_k[\theta(t)]\} \\ \mathcal{L}_k[\Gamma, x : t]\theta &= \{(\varsigma \star \varsigma_x, \gamma[x \mapsto v_x]) \mid (\varsigma, \gamma) \in \mathcal{L}_k[\Gamma]\theta \wedge (\varsigma_x, v_x) \in \mathcal{V}_k[\theta(t)]\} \end{aligned}$$

And so the final semantic interpretation of a typing judgement simply quantifies over all possible fractional permission substitutions θ , linear value substitutions γ , intuitionistic value substitutions δ and heaps σ . Note that, $\varsigma \equiv \mathcal{H}[\theta(\sigma)]$.

$$k[\Theta; \Delta; \Gamma \vdash e : t] = \forall \theta, \delta, \gamma, \sigma. \Theta = \text{dom}(\theta) \wedge (\varsigma, \gamma) \in \mathcal{L}_k[\Gamma]\theta \wedge \delta \in \mathcal{I}_k[\Delta]\theta \Rightarrow (\varsigma, \theta(\delta(\gamma(e)))) \in \mathcal{C}_k[\theta(t)]$$

3.4 Soundness Theorem

Theorem 1 (The Fundamental Lemma of Logical Relations)

$$\forall \Theta, \Delta, \Gamma, e, t. \Theta; \Delta; \Gamma \vdash e : t \Rightarrow \forall k. k[\Theta; \Delta; \Gamma \vdash e : t]$$

To prove the above theorem, we need several lemmas; the interesting ones are: the moral equivalent of the frame rule (C.1), monotonicity for the step-index (C.5), splitting up environments corresponds to splitting up heaps (C.7) and heap-and-expressions take the same steps of evaluation under any substitution of their free fractional permissions (C.8).

The proof proceeds by induction on the typing judgement. The case for `TY_FIX` is the reason we quantify over the step-index k in the *conclusion* of the soundness theorem. It allows us to then induct over the step-index and assume exactly the thing we need to prove at a smaller index.

The case for `TY_GEN` follows a similar pattern, but has the extra complication of reducing an expression with an arbitrary fractional permission variable in it, and then instantiating it at the last moment to conclude, which is where C.8 (heap-and-expressions take the same steps of evaluation under any substitution of their free fractional permissions) is used.

The rest of the cases are either very simple base cases (variables, unit, boolean, integer or element literals) or follow very similar patterns; for these, only `TY_LET` is presented in full and other similar cases simply highlight exactly what would be different. The general idea is to split up the linear substitution and heap along the same split of Γ/Γ' , then (by induction) use $\mathcal{C}_k[-]$ and one ‘half’ of the linear substitution and heap to conclude the ‘first’ sub-expression either takes $j < k$ steps to `err` or another heap-and-expression.

In the first case, you use `OP_CONTEXT_ERR` to conclude the whole let-expression does the same. Similarly we use `OP_CONTEXT` j times in the second case. However, a small book-keeping wrinkle needs to be taken care of in the case that the heap-and-expression turns into a value in $i \leq j$ steps: `OP_CONTEXT` is not functorial for the n -fold iteration of \rightarrow . Basically, the following is not quite true:

$$\frac{\langle \sigma, e \rangle \rightarrow^j \langle \sigma', e' \rangle}{\langle \sigma, C[e] \rangle \rightarrow^j \langle \sigma', C[e'] \rangle} \quad \text{OP_CONTEXT}$$

because after the i steps, we need to invoke `OP_LET_VAR` to proceed evaluation for any remaining $j - i$ steps. After that, it suffices to use the induction hypothesis on the second sub-expression to finish the proof. To do so, we need to construct a valid linear substitution and heap (i.e., one in $\mathcal{L}_k[\Gamma', x : t]\theta$). We take the other ‘half’ of the linear substitution and heap (from the initial split at the start) and extend it with $[x \mapsto v]$, (where x is the variable bound in the let-expression and v is the value we assume the first sub-expression evaluated to in i steps).

4 Implementation

4.1 Implementation Strategy

NUMLIN transpiles to OCaml and its implementation follows the structure of a typical domain-specific language (DSL) compiler. Although NUMLIN’s current implementation is not as an embedded DSL, its the general design is simple enough to adapt to being so and also to target other languages.

Alongside the transpiler, a ‘Read-Check-Translate’ loop, benchmarking program and a test suite are included in the artifacts accompanying this paper.

1. **Parsing.** A generated, LR(1) parser parses a text file into a syntax tree. In general, this part will vary for different languages and can also be dealt with using combinators or syntax-extensions (the EDSL approach) if the host language offers such support.
2. **Desugaring.** The syntax tree is then desugared into a smaller, more concise, abstract syntax tree. This allows for the type checker to be simpler to specify and easier to implement.
3. **Matrix Expressions** are also desugared into the abstract syntax tree through pattern-matching.
4. **Type checking.** The abstract syntax tree is explicitly typed, with some inference to make writing typical programs more convenient.
5. **Code Generation.** The abstract syntax tree is translated into OCaml, with a few ‘optimisations’ to produce more readable code. This process is type-preserving: NUMLIN’s type system is embedded into OCaml’s (Figure 11), so the OCaml type checker acts as a sanity check on the generated code.

A very pleasant way to use NUMLIN is to have the build system generate code at *compile-time* and then have the generated code be used by other modules like normal OCaml functions. This makes it possible and even easy to use NUMLIN alongside existing OCaml libraries; in fact, this is exactly how the benchmarking program and test-suite use code written in NUMLIN.

4.1.1 Desugaring, Matrix Expressions and Type Checking

Desugaring is conventional, outlined in Appendix F. Matrix expressions are translated into BLAS/ LAPACK calls via purely syntactic pattern-matching, outlined in Figure 10.

4.1.2 Type checking

is mostly standard for a linearly typed language, with the exception of fractional permission inference. By restricting fractions to be non-positive integer powers of two, we only need to keep track of the logarithm of the fractions used. Explicit sharing and unsharing removes the need for performing dataflow analysis. As a result, all fractional arithmetic can be solved with unification, and in doing so, fractions become directly usable in NUMLIN's type-system as opposed to a convenient theoretical tool.

Because all functions must have their argument types explicitly annotated, inferring the correct fraction at a call-site is simply a matter of unification. We believe *full-inference of fractional permissions is similarly just matter of unification* (thanks to an experimental implementation of just this feature), even though the formal system we present here is for an explicitly-typed language.

There are a few differences between the type system as presented in 3.2 and how we implemented it: the environment *changes* as a result of type checking an expression (the standard transformation to avoid a non-deterministic split of the environment for checking pairs); variables are *marked as used* rather than removed for better error messages; variables are *tagged* as linear or intuitionistic in *one* environment as opposed to being stored in *two* separate ones (this allows scoping/variable look-up to be handled uniformly).

4.1.3 Code Generation

is a straightforward mapping from NUMLIN's core constructs to high-level OCaml ones. We embed NUMLIN's type- and term- constructors into OCaml as a sanity check on the output (Figure 11).

This is also useful when using NUMLIN from within OCaml; for example, we can use existing tools to inspect the type of the function we are using (Figure 12). It is worth reiterating that only the type- and term- constructors are translated into OCaml, NUMLIN's precise control over linearity and aliasing are not brought over.

We actually use this fact to our advantage to clean up the output OCaml by removing what would otherwise be redundant re-bindings (Figure 13). Combined with a code-formatter, the resulting code is not obviously correct and exactly what an expert would intend to write by hand, but now with the guarantees and safety of NUMLIN behind it. A small example is shown in Figure 14, a larger one in Figure 16.

4.2 Performance Metrics

We think that using NUMLIN has two primary benefits: safety and performance. We discuss safety in 5.1, where we describe how we used NUMLIN to find linearity and aliasing bugs in a linear algebra algorithm that was *generated* by another program.

4.2.1 Setup

For performance, we measured the execution times of four equivalent implementations of a Kalman filter: in C (using CBLAS), NUMLIN (using OWL's low-level CBLAS bindings), OCaml

$\text{let } v \leftarrow x[e] \text{ in } e \Rightarrow \text{let } (x, !v) = x[e] \text{ in } e$ (similarly for matrices)
 $\text{let } x_2 \leftarrow \text{new } [|x_1|] \text{ in } e \Rightarrow \text{let } (x_1, x_2) = \text{copyM_} x_1 \text{ in } e$
 $\text{let } x_2 \leftarrow [|x_1|] \text{ in } e \Rightarrow \text{let } (x_1, x_2) = \text{copyM_to_} x_1 \text{ in } e$
 $M ::= X \mid X^T \mid \text{sym}(X)$

 $\text{let } Y \leftarrow \text{new } (n, k) [| \alpha M_1 M_2 |] \text{ in } e \Rightarrow$
 $\quad \text{let } Y = \text{matrix } n \text{ } k \text{ in let } Y \leftarrow [| \alpha M_1 M_2 + 0Y |] \text{ in } e$
 $\text{let } Y \leftarrow [| \alpha X X^T + \beta Y |] \text{ in } e \Rightarrow$
 $\quad \text{let } (X, Y) = \text{syrk false } \alpha _ X \beta Y \text{ in } e$
 $\text{let } Y \leftarrow [| \alpha X^T X + \beta Y |] \text{ in } e \Rightarrow$
 $\quad \text{let } (X, Y) = \text{syrk true } \alpha _ X \beta Y \text{ in } e$
 $\text{let } Y \leftarrow [| \alpha \text{sym}(X_1) X_2 + \beta Y |] \text{ in } e \Rightarrow$
 $\quad \text{let } ((X_1, X_2), Y) = \text{symm false } \alpha _ X_1 _ X_2 \beta Y \text{ in } e$
 $\text{let } Y \leftarrow [| \alpha X_2 \text{sym}(X_1) + \beta Y |] \text{ in } e \Rightarrow$
 $\quad \text{let } ((X_1, X_2), Y) = \text{symm true } \alpha _ X_1 _ X_2 \beta Y \text{ in } e$
 $\text{let } Y \leftarrow [| \alpha X_1^{T?} X_2^{T?} + \beta Y |] \text{ in } e \Rightarrow$
 $\quad \text{let } ((X_1, X_2), Y) = \text{gemm } \alpha _ (X_1, \text{true}_{\text{false}}) _ (X_2, \text{true}_{\text{false}}) \beta Y \text{ in } e$

■ **Figure 10** Purely syntactic pattern-matching translations of matrix expressions.

518 (using OWL’s intended, safe/copying-by-default interface), and Python (using NUMPY, with
 519 the interpreter started and functions interpreted). We measured execution time in micro-
 520 seconds, against an exponentially (powers of 5) increasing scaling factor for matrix size
 521 parameters $n = 5$ and $k = 3$.

522 For large scaling factors ($n = 5^4, 5^5$), we triggered a full garbage-collection before
 523 measuring the execution time of a single call of a function. However, due to the limitations
 524 of the micro-benchmarking library we used, for smaller scaling factors ($n = 5^1, 5^2, 5^3$), we
 525 measured the execution time of *multiple* calls to a function in a loop, thus including potential
 526 garbage-collection effects.

527 We also measured the execution times of L1-norm minimisation and the “linear-regression”
 528 $((\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y})$ similarly, but without a C implementation.

529 4.2.2 Hypothesis

530 We expected the C implementation to be faster than the NUMLIN one because the latter has
 531 the additional (but relatively low) overhead of dimension checks and crossing the OCaml/C
 532 FFI for each call to a CBLAS routine, even though the calls and their order are exactly the
 533 same. We expected the OCaml and Python implementations to be slower because they
 534 allocate more temporaries (so possibly less cache-friendly) and carry out more floating-point
 535 operations – the CBLAS and NUMLIN implementations use ternary kernels (coalescing steps),
 536 a Cholesky decomposition (of a symmetric matrix, which is more efficient than the LU

$f ::=$		
fc	<code>module Arr =</code>	$\llbracket fc \rrbracket = 'fc$
Z	<code>Owl.Dense.Ndarray.D</code>	$\llbracket Z \rrbracket = z$
$S f$	<code>type z = Z</code>	$\llbracket S f \rrbracket = \llbracket f \rrbracket s$
	<code>type 'a s = Succ</code>	$\llbracket unit \rrbracket = unit$
$t ::=$		$\llbracket bool \rrbracket = bool$
<code>unit</code>	<code>type 'a arr =</code>	$\llbracket int \rrbracket = int$
<code>bool</code>	<code>A of Arr.arr</code>	$\llbracket elt \rrbracket = float$
<code>int</code>	<code>[@@unboxed]</code>	$\llbracket f arr \rrbracket = \llbracket f \rrbracket arr$
<code>elt</code>		$\llbracket f mat \rrbracket = \llbracket f \rrbracket mat$
<code>f arr</code>	<code>type 'a mat =</code>	$\llbracket ! t \rrbracket = \llbracket t \rrbracket bang$
<code>f mat</code>	<code>M of Arr.arr</code>	$\llbracket \forall fc. t \rrbracket = \llbracket t \rrbracket$
<code>! t</code>	<code>[@@unboxed]</code>	$\llbracket t \otimes t' \rrbracket = \llbracket t \rrbracket * \llbracket t' \rrbracket$
$\forall fc. t$	<code>type 'a bang =</code>	$\llbracket t \multimap t' \rrbracket = \llbracket t \rrbracket \rightarrow \llbracket t' \rrbracket$
$t \otimes t'$	<code>Many of 'a</code>	
$t \multimap t'$	<code>[@@unboxed]</code>	

■ **Figure 11** NUMLIN's type grammar (left) and its embedding into OCaml (right).

```

1 let lt4la_kalman ~sigma ~h ~mu ~r ~data =
0   Examples.Kalman.it (M sigma) (M h) (M mu) (M r) (M data)
NORMAL test/examples_test.ml
'a mat ->
'b mat ->
'c mat ->
z mat ->
z mat -> ('a mat * ('b mat * ('c mat * (z mat * z mat)))) * (z mat * z mat)
:merlin-type-history:
0   let fact = Examples.Factorial.it in
NORMAL test/examples_test.ml
int bang -> int bang

```

■ **Figure 12** Using NUMLIN functions from OCaml.

537 decomposition used for inverting a matrix in OWL and NUMPY) and `symm` (symmetric matrix
538 multiplication, halving the number of floating-point multiplications required).

539 4.2.3 Results

540 The results in Figures 15 are as we expected: C is the fastest, followed by NUMLIN, with
541 OCaml and Python last. Differences in timings are quite pronounced at small matrix sizes,
542 but are still significant at larger ones. Specifically for the Kalman filter, for $n = 625$, CBLAS
543 took $112 \pm 35 ms$, NUMLIN took $105 \pm 25 ms$, OWL took $124 \pm 38 ms$ and NUMPY took
544 $112 \pm 12 ms$; for $n = 3125$, CBLAS took $10.8 \pm 0.7 s$, NUMLIN took $12.0 \pm 1.2 s$, OWL took
545 $13.3 \pm 0.2 s$ and NUMPY took $12.7 \pm 0.6 s$.

546 Worth highlighting here is the other major advantage of using NUMLIN is reduced
547 memory usage. Whilst the OWL and NUMPY use 11 temporary matrices for the Kalman
548 filter, (excluding the 2 matrices which store the results), using $n + n^2 + 4nk + 3k^2 + 2k \approx 4n^2$
549 (for $k = 3n/5$) words of memory, CBLAS and NUMLIN use only 2 temporary matrices
550 (excluding the one matrix which stores one of the results), using only $n^2 + nk \leq 2n^2$ words

```

let Many x = x in
let Many x = Many (Many x) in <exp>  ⇒  <exp>

(* fixp = fix (f, x:t, <exp> : t') *)
(*1*) let Many f = Many fixp in <body>  ⇒  let rec f x = <exp> in <body>
(*2*) let f = fixp in <body>

(*1*) let Many x = <exp> in
(*-*) let Many x = Many (Many x) in <body>  ⇒  let x = <exp> in <body>
(*2*) let Many x = Many <exp> in <body>
(*3*) (fun x : t -> <body>) <exp>

```

■ **Figure 13** Removing redundant re-bindings during translation to OCaml.

```

let rec f i n x0 row =
  if Prim.extract @@ Prim.eqI i n then (row, x0)
  else
    let row, x1 = Prim.get row i in
    f (Prim.addI i (Many 1)) n (Prim.addE x0 x1) row
in
f

```

■ **Figure 14** Recursive OCaml function for a summing over an array, generated (at *compile time*) from the code in Figure 2, passed through `ocamlformat` for presentation.

551 of memory.

552 4.2.4 Analysis

553 As matrix sizes increase, assuming sufficient memory, the difference in the number of floating-
 554 point operations ($O(n^3)$) dominates execution times. However for small matrix sizes, since
 555 n is small and the measurements were over multiple calls to a function in a loop, the large
 556 number of temporaries show the adverse effect of not re-using memory at even quite small
 557 matrix sizes: creating pressure on the garbage collector.

558 5 Discussion and Related Work

559 5.1 Finding Bugs in SymPy's Output

560 Prior to this project, we had little experience with linear algebra libraries or the problem
 561 of matrix expression compilation. As such, we based our initial NUMLIN implementation
 562 of a Kalman filter using BLAS and LAPACK, on a popular GitHub gist of a Fortran
 563 implementation, one that was *automatically generated* from SymPy's matrix expression
 564 compiler [15].

565 Once we translated the implementation from Fortran to NUMLIN, we attempted to compile
 566 it and found that (to our surprise) it did not type-check. This was because the original
 567 implementation contained incorrect aliasing, unused variables and unnecessary temporaries,
 568 and did not adhere to Fortran's read/write permissions (with respect to `intent` annotations
 569 `in`, `out` and `inout`) all of which were now highlighted by NUMLIN's type system.

570 The original implementation used 6 temporaries, one of which was immediately spotted
 571 as never being used due to linearity. It also contained two variables which were marked as

`intent(in)` but would have been written over by calls to ‘`gemm`’, spotted by the fractional-capabilities feature. Furthermore, it used a matrix *twice* in a call to ‘`symm`’, once with a read permission but once with a *write* permission. Fortran assumes that any parameter being written to is not aliased and so this call was not only incorrect, but illegal according to the standard, both aspects of which were captured by linearity and fractional-capabilities.

Lastly, it contained another unnecessary temporary, however one that was not obvious without linear types. To spot it, we first performed live-range splitting (checked by linearity) by hoisting calls to `freeM` and then annotated the freed matrices with their dimensions. After doing so and spotting two disjoint live-ranges of the same size, we replaced a call to `freeM` followed by allocating call to `copy` with one, in-place call to `copyM_to`. We believe the ability to boldly refactor code which manages memory is good evidence of the usefulness of linearity as a tool for programming.

5.2 Related Work

Using linear types for BLAS routines is a particularly good domain fit (given the implicit restrictions on aliasing arguments), and as a result the idea of using substructural types to express array computations is not a particularly new one [16, 10, 4]. However, many of these designs have been focused on building languages to *implement* the kernel linear algebra functions, and as a result, they tend to add additional limitations on the language design. Both Futhark [10] and Single Assignment C [16] omit higher-order functions facilitating compilation to GPUs. The work of [4] forbids term-level recursion, in order to ensure that all higher-order computations can be statically normalized away and thereby maximize opportunities for array fusion.

In contrast, our approach is to begin with the assumption that we can take existing efficient BLAS-like libraries, and then enforce their correct *usage* using a linear type discipline with fractional permissions.

This approach is similar to the one taken in linear algebra libraries for Rust – these libraries typically take advantage of the distinction that Rust’s type system offers between mutable views/references to arrays. The work of [18] and [11] suggest that Rust’s borrow-checker can be expressed in simpler terms using *fractional-permissions*, though to our knowledge the programmer-visible lifetime analysis in Rust has never been formalized.

Using fractional permissions directly, as we did, rather than relying on a complex has two main benefits. First, our type system demonstrates that type systems for fractional permissions can be *dramatically* simpler than existing state-of-the-art approaches, including both industrial languages like Rust, as well as the academic languages developed by [5]. Their type system, much like Rust, build a complex dataflow analysis into the typing rules to infer when variables can be shared or not. This allows for very natural-looking user programs, but creates the impression that using fractional permissions requires a heavy theoretical and engineering effort going well beyond that needed for supporting basic linear types.

Instead, our approach of requiring sharing to be made explicit means that the existing unification machinery already in place for ordinary ML-style type inference, can be reused to support fractions. Basically, we can view sharing a value as dividing a fraction by two, and after taking logarithms all fractions are Peano numbers, whose equality can be established with ordinary unification.

This fact is important because there are major upcoming implementations of linear types such as Linear Haskell [3], which do not have built-in support for fractional permissions. Instead, Linear Haskell takes a slightly different definition of linearity, one based on *arrows* as opposed to *kinds*: for $f : a \multimap b$, if fu is used exactly once *then* u is used exactly once. Whilst

this has the advantage of being backwards-compatible, it also means that the type system has no built-in support for the concurrent reader, exclusive writer pattern that fractional permissions enable.

However, since our type system shows demonstrates unification is “all one needs” for fractions, it should be possible to *encode* NUMLIN’s approach to fractional permissions in Linear Haskell by adding a GADT-style natural number index to array types tracking the fraction, which should enable supporting high-performance BLAS bindings in Linear Haskell. Actually implementing this is something we leave for future work, as there remains an issue which we do not see a good encoding for. Namely, only having support for linear functions makes it a bit inconvenient to manipulate linear values directly – programs end up taking on a CPS-like structure. This seems to remain an advantage of a direct implementation of linear types.

5.3 Simplicity and Further Work

We are pleasantly surprised at how simple the overall design and implementation of NUMLIN is, given its expressive power and usability. So simple in fact, that fractions, a convenient theoretical abstraction until this point, could be implemented by restricting division and multiplication to be by 2 only [7], thus turning any required arithmetic into unification.

Indeed, the focus on getting a working prototype early on (so that we could test it with real BLAS/LAPACK routines as soon as possible) meant that we only added features to the type system when it was clear that they were absolutely necessary: these features were !-types and value-restriction for the [Many](#) constructor.

Going forwards, one may wish to eliminate even more runtime errors from NUMLIN, by extending its type system. For example, we could have used existential types to statically track pointer identities[14], or parametric polymorphism.

We could also attempt to catch mismatched dimensions at compile time as well. While this could be done with generative phantom types[1], using dependent types may offer more flexibility in *partitioning* regions[13] or statically enforcing dimensions related constraints of the arguments at compile-time. ATS[9] is just such a language which combines dependent and linear types; although it provides BLAS bindings, it does not aim to provide aliasing restrictions as demonstrated in this paper.

Taking this idea one step even further, since matrix dimensions are typically fixed at runtime, we could *stage* NUMLIN programs and compile matrix expressions using more sophisticated algorithms[2]. However, it is worth noting that without care, such algorithms[15], usually based on graph-based, ad-hoc dataflow analysis, can produce erroneous output which would not get past a linear type system with fractions.

We also think that this concept (and the general design of its implementation) need not be limited to linear algebra: we could conceivably ‘backport’ this idea to other contexts that need linearity (concurrency, single-use continuations, zero-copy buffer, streaming I/O) or combine it with dependent types to achieve even more expressive power to split up a single block of memory into multiple regions in an arbitrary manner[13].

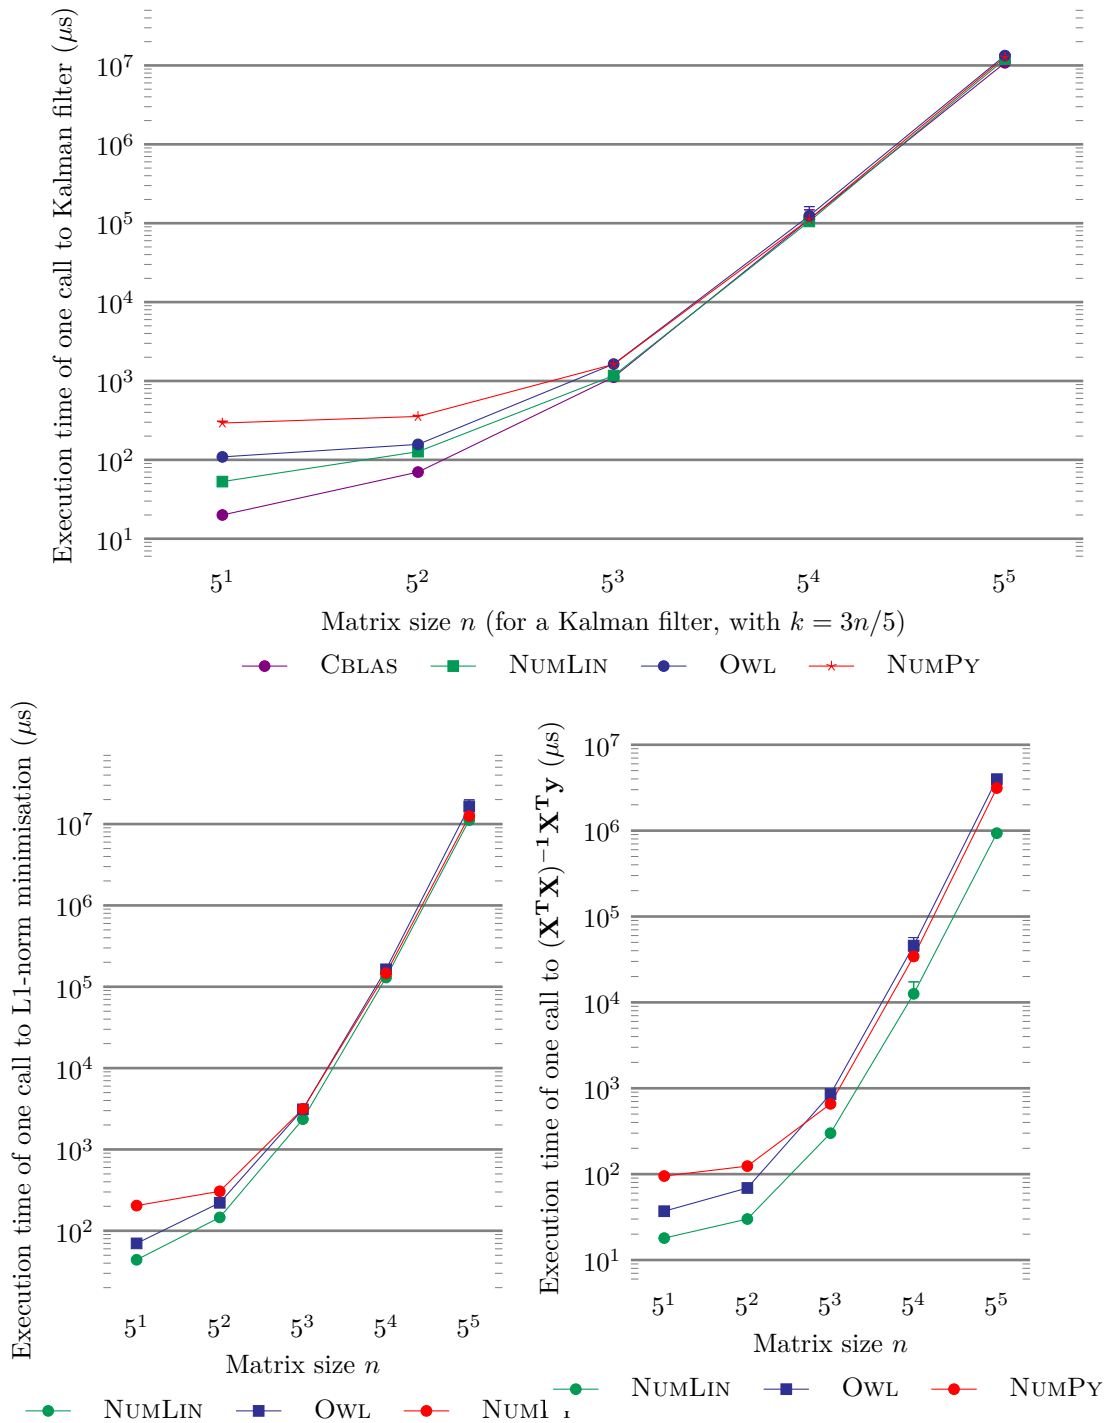


Figure 15 Comparison of execution times (error bars are present but quite small). Small matrices and timings $n \leq 5^3$ were micro-benchmarked with the Core_bench library. Larger ones used Unix's `getrusage` functionality, sandwiched between calls to `Gc.full_major` for the OCaml implementations.

References

- 1 Akinori Abe and Eijiro Sumii. A simple and practical linear algebra library interface with static size checking. *arXiv preprint arXiv:1512.01898*, 2015.
- 2 Henrik Barthels, Marcin Copik, and Paolo Bientinesi. The generalized matrix chain algorithm. *arXiv preprint arXiv:1804.04021*, 2018.
- 3 Jean-Philippe Bernardy, Mathieu Boespflug, Ryan R Newton, Simon Peyton Jones, and Arnaud Spiwack. Linear haskell: practical linearity in a higher-order polymorphic language. *Proceedings of the ACM on Programming Languages*, 2(POPL):5, 2017.
- 4 Jean-Philippe Bernardy, Victor López Juan, and Josef Svenningsson. Composable efficient array computations using linear types. *Unpublished Draft*, 2016.
- 5 Kevin Bierhoff, Nels E Beckman, and Jonathan Aldrich. Fraction polymorphic permission inference.
- 6 Richard Bornat, Cristiano Calcagno, Peter O’Hearn, and Matthew Parkinson. Permission accounting in separation logic. In *ACM SIGPLAN Notices*, volume 40, pages 259–270. ACM, 2005.
- 7 John Boyland. Checking interference with fractional permissions. In *International Static Analysis Symposium*, pages 55–72. Springer, 2003.
- 8 Alex Bronstein, Yoni Choukroun, Ron Kimmel, and Matan Sela. Consistent discretization and minimization of the l1 norm on manifolds. In *3D Vision (3DV), 2016 Fourth International Conference on*, pages 435–440. IEEE, 2016.
- 9 Sa Cui, Kevin Donnelly, and Hongwei Xi. Ats: A language that combines programming with theorem proving. In *International Workshop on Frontiers of Combining Systems*, pages 310–320. Springer, 2005.
- 10 Troels Henriksen, Niels GW Serup, Martin Elsmann, Fritz Henglein, and Cosmin E Oancea. Futhark: purely functional gpu-programming with nested parallelism and in-place array updates. *ACM SIGPLAN Notices*, 52(6):556–571, 2017.
- 11 Ralf Jung, Jacques-Henri Jourdan, Robbert Krebbers, and Derek Dreyer. Rustbelt: securing the foundations of the rust programming language. *PACMPL*, 2(POPL):66:1–66:34, 2018. URL: <https://doi.org/10.1145/3158154>, doi:10.1145/3158154.
- 12 Rudolph Emil Kalman. A new approach to linear filtering and prediction problems. *Journal of basic Engineering*, 82(1):35–45, 1960.
- 13 Conor McBride. Code mesh london 2016, keynote: Spacemonads. <https://www.youtube.com/watch?v=QoJLQY5H0RI>. Accessed: 08/05/2018.
- 14 Greg Morrisett, Amal Ahmed, and Matthew Fluet. L 3: a linear language with locations. In *International Conference on Typed Lambda Calculi and Applications*, pages 293–307. Springer, 2005.
- 15 Matthew Rocklin. *Mathematically informed linear algebra codes through term rewriting*. PhD thesis, 2013.
- 16 Sven-Bodo Scholz. Single assignment c: efficient support for high-level array operations in a functional setting. *Journal of functional programming*, 13(6):1005–1059, 2003.
- 17 Philip Wadler. Linear types can change the world. In M. Broy and C. B. Jones, editors, *IFIP TC 2 Working Conference on Programming Concepts and Methods*, pages 561–581, Sea of Gallilee, Israel, April 1990. North-Holland.
- 18 Aaron Weiss, Daniel Patterson, and Amal Ahmed. Rust distilled: An expressive tower of languages. *arXiv preprint arXiv:1806.02693*, 2018.

704 **A** NumLin Specification

705 **A.1** Static Semantics

706 $\boxed{\Theta; \Delta; \Gamma \vdash e : t}$ Typing rules for expressions

707 $\frac{}{\Theta; \Delta; \cdot, x : t \vdash x : t} \text{TY_VAR_LIN}$

708 $\frac{x : t \in \Delta}{\Theta; \Delta; \cdot \vdash x : t} \text{TY_VAR}$

709 $\frac{\begin{array}{l} \Theta; \Delta; \Gamma \vdash e : t \\ \Theta; \Delta; \Gamma', x : t \vdash e' : t' \end{array}}{\Theta; \Delta; \Gamma, \Gamma' \vdash \text{let } x = e \text{ in } e' : t'} \text{TY_LET}$

710 $\frac{}{\Theta; \Delta; \cdot \vdash () : \text{unit}} \text{TY_UNIT_INTRO}$

711 $\frac{\begin{array}{l} \Theta; \Delta; \Gamma \vdash e : \text{unit} \\ \Theta; \Delta; \Gamma' \vdash e' : t \end{array}}{\Theta; \Delta; \Gamma, \Gamma' \vdash \text{let } () = e \text{ in } e' : t} \text{TY_UNIT_ELIM}$

712 $\frac{}{\Theta; \Delta; \cdot \vdash \text{true} : \text{bool}} \text{TY_BOOL_TRUE}$

713 $\frac{}{\Theta; \Delta; \cdot \vdash \text{false} : \text{bool}} \text{TY_BOOL_FALSE}$

714 $\frac{\begin{array}{l} \Theta; \Delta; \Gamma \vdash e : !\text{bool} \\ \Theta; \Delta; \Gamma' \vdash e_1 : t' \\ \Theta; \Delta; \Gamma' \vdash e_2 : t' \end{array}}{\Theta; \Delta; \Gamma, \Gamma' \vdash \text{if } e \text{ then } e_1 \text{ else } e_2 : t} \text{TY_BOOL_ELIM}$

715 $\frac{}{\Theta; \Delta; \cdot \vdash k : \text{int}} \text{TY_INT_INTRO}$

716 $\frac{}{\Theta; \Delta; \cdot \vdash el : \text{elt}} \text{TY_ELT_INTRO}$

717 $\frac{\begin{array}{l} \Theta; \Delta; \cdot \vdash v : t \\ v \neq l \end{array}}{\Theta; \Delta; \cdot \vdash \text{Many } v : !t} \text{TY_BANG_INTRO}$

718 $\frac{\begin{array}{l} \Theta; \Delta; \Gamma \vdash e : !t \\ \Theta; \Delta, x : t; \Gamma' \vdash e' : t' \end{array}}{\Theta; \Delta; \Gamma, \Gamma' \vdash \text{let Many } x = e \text{ in } e' : t'} \text{TY_BANG_ELIM}$

719 $\frac{\begin{array}{l} \Theta; \Delta; \Gamma \vdash e : t \\ \Theta; \Delta; \Gamma' \vdash e' : t' \end{array}}{\Theta; \Delta; \Gamma, \Gamma' \vdash (e, e') : t \otimes t'} \text{TY_PAIR_INTRO}$

23:24 NumLin: Linear Types for Linear Algebra

$$\frac{\begin{array}{c} \Theta; \Delta; \Gamma \vdash e_{12} : t_1 \otimes t_2 \\ \Theta; \Delta; \Gamma', a : t_1, b : t_2 \vdash e : t \end{array}}{\Theta; \Delta; \Gamma, \Gamma' \vdash \mathbf{let} (a, b) = e_{12} \mathbf{in} e : t} \quad \text{TY_PAIR_ELIM}$$

$$\frac{\begin{array}{c} \Theta \vdash t' \text{Type} \\ \Theta; \Delta; \Gamma, x : t' \vdash e : t \end{array}}{\Theta; \Delta; \Gamma \vdash \mathbf{fun} x : t' \rightarrow e : t' \multimap t} \quad \text{TY_LAMBDA}$$

$$\frac{\begin{array}{c} \Theta; \Delta; \Gamma \vdash e : t' \multimap t \\ \Theta; \Delta; \Gamma' \vdash e' : t' \end{array}}{\Theta; \Delta; \Gamma, \Gamma' \vdash e e' : t} \quad \text{TY_APP}$$

$$\frac{\Theta, fc; \Delta; \Gamma \vdash e : t}{\Theta; \Delta; \Gamma \vdash \mathbf{fun} fc \rightarrow e : \forall fc. t} \quad \text{TY_GEN}$$

$$\frac{\begin{array}{c} \Theta \vdash f \text{Perm} \\ \Theta; \Delta; \Gamma \vdash e : \forall fc. t \end{array}}{\Theta; \Delta; \Gamma \vdash e[f] : t[f/fc]} \quad \text{TY_SPC}$$

$$\frac{\Theta; \Delta, g : t \multimap t'; \cdot, x : t \vdash e : t'}{\Theta; \Delta; \cdot \vdash \mathbf{fix} (g, x : t, e : t') : t \multimap t'} \quad \text{TY_FIX}$$

726

727 A.2 Dynamic Semantics

728 $\langle \sigma, e \rangle \rightarrow \text{Config}$ Operational semantics

$$\frac{}{\langle \sigma, \mathbf{let} () = () \mathbf{in} e \rangle \rightarrow \langle \sigma, e \rangle} \quad \text{OP_LET_UNIT}$$

$$\frac{}{\langle \sigma, \mathbf{let} x = v \mathbf{in} e \rangle \rightarrow \langle \sigma, e[x/v] \rangle} \quad \text{OP_LET_VAR}$$

$$\frac{}{\langle \sigma, \mathbf{if} (\mathbf{Many true}) \mathbf{then} e_1 \mathbf{else} e_2 \rangle \rightarrow \langle \sigma, e_1 \rangle} \quad \text{OP_IF_TRUE}$$

$$\frac{}{\langle \sigma, \mathbf{if} (\mathbf{Many false}) \mathbf{then} e_1 \mathbf{else} e_2 \rangle \rightarrow \langle \sigma, e_2 \rangle} \quad \text{OP_IF_FALSE}$$

$$\frac{}{\langle \sigma, \mathbf{let Many} x = \mathbf{Many} v \mathbf{in} e \rangle \rightarrow \langle \sigma, e[x/v] \rangle} \quad \text{OP_LET_MANY}$$

$$\frac{}{\langle \sigma, \mathbf{let} (a, b) = (v_1, v_2) \mathbf{in} e \rangle \rightarrow \langle \sigma, e[a/v_1][b/v_2] \rangle} \quad \text{OP_LET_PAIR}$$

$$\frac{}{\langle \sigma, (\mathbf{fun} fc \rightarrow v)[f] \rangle \rightarrow \langle \sigma[f c/f], v[f c/f] \rangle} \quad \text{OP_FRAC_PERM}$$

$$\frac{}{\langle \sigma, \mathbf{fix} (g, x : t, e : t') v \rangle \rightarrow \langle \sigma, e[x/v][g/\mathbf{fix} (g, x : t, e : t')] \rangle} \quad \text{OP_APP_FIX}$$

$$\frac{}{\langle \sigma, (\mathbf{fun} x : t \rightarrow e) v \rangle \rightarrow \langle \sigma, e[x/v] \rangle} \quad \text{OP_APP_LAMBDA}$$

737

738	$\frac{\langle \sigma, e \rangle \rightarrow \langle \sigma', e' \rangle}{\langle \sigma, C[e] \rangle \rightarrow \langle \sigma', C[e'] \rangle} \quad \text{OP_CONTEXT}$	
739	$\frac{\langle \sigma, e \rangle \rightarrow \mathbf{err}}{\langle \sigma, C[e] \rangle \rightarrow \mathbf{err}} \quad \text{OP_CONTEXT_ERR}$	
740	$\frac{0 \leq k_1, k_2 \quad l \text{ fresh}}{\langle \sigma, \mathbf{matrix} \ k_1 \ k_2 \rangle \rightarrow \langle \sigma + \{l \mapsto_1 M_{k_1, k_2}\}, l \rangle} \quad \text{OP_MATRIX}$	
741	$\frac{k_1 < 0 \text{ or } k_2 < 0}{\langle \sigma, \mathbf{matrix} \ k_1 \ k_2 \rangle \rightarrow \mathbf{err}} \quad \text{OP_MATRIX_NEG}$	
742	$\frac{}{\langle \sigma + \{l \mapsto_1 m_{k_1, k_2}\}, \mathbf{free} \ l \rangle \rightarrow \langle \sigma, () \rangle} \quad \text{OP_FREE}$	
743	$\frac{}{\langle \sigma + \{l \mapsto_f m_{k_1, k_2}\}, \mathbf{share}[f] \ l \rangle \rightarrow \langle \sigma + \{l \mapsto_{\frac{1}{2}f} m_{k_1, k_2}\} + \{l \mapsto_{\frac{1}{2}f} m_{k_1, k_2}\}, (l, l) \rangle} \quad \text{OP_SHARE}$	
744	$\frac{}{\langle \sigma + \{l \mapsto_{\frac{1}{2}f} m_{k_1, k_2}\} + \{l \mapsto_{\frac{1}{2}f} m_{k_1, k_2}\}, \mathbf{unshare}[f] \ l \ l \rangle \rightarrow \langle \sigma + \{l \mapsto_f m_{k_1, k_2}\}, l \rangle} \quad \text{OP_UNSHARE_EQ}$	
745	$\frac{l \neq l'}{\langle \sigma + \{l \mapsto_{\frac{1}{2}f} m_{k_1, k_2}\} + \{l' \mapsto_{\frac{1}{2}f} m'_{k_1, k_2}\}, \mathbf{unshare}[f] \ l \ l' \rangle \rightarrow \mathbf{err}} \quad \text{OP_UNSHARE_NEQ}$	
746	$\frac{\begin{array}{l} \sigma' \equiv \sigma + \{l_1 \mapsto_{fc_1} m_{1 \ k_1, k_2}\} + \{l_2 \mapsto_{fc_2} m_{2 \ k_2, k_3}\} \\ \sigma_1 \equiv \sigma' + \{l_3 \mapsto_1 m_{3 \ k_1, k_3}\} \\ \sigma_2 \equiv \sigma' + \{l_3 \mapsto_1 (m_1 \ m_2 + m_3)_{k_1, k_3}\} \end{array}}{\langle \sigma_1, \mathbf{gemm}[fc_1] \ l_1[fc_2] \ l_2 \ l_3 \rangle \rightarrow \langle \sigma_2, ((l_1, l_2), l_3) \rangle} \quad \text{OP_GEMM_MATCH}$	
747	$\frac{k_2 \neq k'_2 \quad \sigma' \equiv \sigma + \{l_1 \mapsto_{fc_1} m_{1 \ k_1, k_2}\} + \{l_2 \mapsto_{fc_2} m_{2 \ k'_2, k_3}\}}{\langle \sigma' + \{l_3 \mapsto_1 m_{1 \ k_1, k_3}\}, \mathbf{gemm}[fc_1] \ l_1[fc_2] \ l_2 \ l_3 \rangle \rightarrow \mathbf{err}} \quad \text{OP_GEMM_MISMATCH}$	
748		

749 **B** Interpretation

750 **B.1** Definitions

751 Operationally, $Heap \sqsubseteq Loc \times Permission \times Matrix$ (a multiset), denoted with a σ .
 752 Define its *interpretation* to be $Loc \rightarrow Permission \times Matrix$ with $\star : Heap \times Heap \rightarrow Heap$ as
 753 follows:

$$754 \quad (\varsigma_1 \star \varsigma_2)(l) \equiv \begin{cases} \varsigma_1(l) & \text{if } l \in \text{dom}(\varsigma_1) \wedge l \notin \text{dom}(\varsigma_2) \\ \varsigma_2(l) & \text{if } l \in \text{dom}(\varsigma_2) \wedge l \notin \text{dom}(\varsigma_1) \\ (f_1 + f_2, m) & \text{if } (f_1, m) = \varsigma_1(l) \wedge (f_2, m) = \varsigma_2(l) \wedge f_1 + f_2 \leq 1 \\ \text{undefined} & \text{otherwise} \end{cases}$$

755 Commutativity and associativity of \star follows from that of $+$.
 756 $\varsigma_1 \star \varsigma_2$ is *defined* if it is for all $l \in \text{dom}(\varsigma_1) \cup \text{dom}(\varsigma_2)$.
 757 Define $\mathcal{H}[\![\sigma]\!] = \star_{(l,f,m) \in \sigma} [l \mapsto_f m]$ and **implicitly denote** $\varsigma \equiv \mathcal{H}[\![\theta(\sigma)]\!]$.
 758
 759 The n -fold iteration for the \rightarrow (functional) relation, is also a (functional) relation:

$$760 \quad \forall n. \mathbf{err} \rightarrow^n \mathbf{err} \quad \langle \sigma, v \rangle \rightarrow^n \langle \sigma, v \rangle \quad \langle \sigma, e \rangle \rightarrow^0 \langle \sigma, e \rangle \quad \langle \sigma, e \rangle \rightarrow^{n+1} ((\langle \sigma, e \rangle \rightarrow) \rightarrow^n)$$

761

762 Hence, all bounded iterations end in either an **err**, a heap-and-expression or a heap-and-value.

B.2 Interpretation

$$\begin{aligned}
\mathcal{V}_k[\mathbf{unit}] &= \{(\emptyset, *)\} \\
\mathcal{V}_k[\mathbf{bool}] &= \{(\emptyset, true), (\emptyset, false)\} \\
\mathcal{V}_k[\mathbf{int}] &= \{(\emptyset, n) \mid 2^{-63} \leq n \leq 2^{63} - 1\} \\
\mathcal{V}_k[\mathbf{elt}] &= \{(\emptyset, f) \mid f \text{ a IEEE Float64}\} \\
\mathcal{V}_k[f \mathbf{mat}] &= \{(\{l \mapsto_{2-f} _ \}, l)\} \\
\mathcal{V}_k[!t] &= \{(\emptyset, \mathbf{Many} \ v) \mid (\emptyset, v) \in \mathcal{V}_k[t]\} \\
\mathcal{V}_k[\forall fc. t] &= \{(\varsigma, \mathbf{fun} \ fc \rightarrow v) \mid \forall f. (\varsigma[fc/f], v[fc/f]) \in \mathcal{V}_{k-1}[t[fc/f]]\} \\
\mathcal{V}_k[t_1 \otimes t_2] &= \{(\varsigma_1 \star \varsigma_2, (v_1, v_2)) \mid (\varsigma_1, v_1) \in \mathcal{V}_k[t_1] \wedge (\varsigma_2, v_2) \in \mathcal{V}_k[t_2]\} \\
\mathcal{V}_k[t' \multimap t] &= \{(\varsigma_v, v) \mid (v \equiv \mathbf{fun} \ x : t' \rightarrow e \vee v \equiv \mathbf{fix}(g, x : t', e : t)) \wedge \\
&\quad \forall j \leq k, (\varsigma_{v'}, v') \in \mathcal{V}_j[t']. \varsigma_v \star \varsigma_{v'} \text{ defined} \Rightarrow (\varsigma_v \star \varsigma_{v'}, v \ v') \in \mathcal{C}_j[t]\} \\
\mathcal{C}_k[t] &= \{(\varsigma_s, e_s) \mid \forall j < k, \sigma_r. \varsigma_s \star \sigma_r \text{ defined} \Rightarrow \langle \sigma_s + \sigma_r, e_s \rangle \rightarrow^j \mathbf{err} \vee \exists \sigma_f, e_f. \\
&\quad \langle \sigma_s + \sigma_r, e_s \rangle \rightarrow^j \langle \sigma_f + \sigma_r, e_f \rangle \wedge (e_f \text{ is a value} \Rightarrow (\varsigma_f \star \sigma_r, e_f) \in \mathcal{V}_{k-j}[t])\} \\
\mathcal{I}_k[\cdot]\theta &= \{\emptyset\} \\
\mathcal{I}_k[\Delta, x : t]\theta &= \{\delta[x \mapsto v_x] \mid \delta \in \mathcal{I}_k[\Delta]\theta \wedge (\emptyset, v_x) \in \mathcal{V}_k[\theta(t)]\} \\
\mathcal{L}_k[\cdot]\theta &= \{(\emptyset, [])\} \\
\mathcal{L}_k[\Gamma, x : t]\theta &= \{(\varsigma \star \varsigma_x, \gamma[x \mapsto v_x]) \mid (\varsigma, \gamma) \in \mathcal{L}_k[\Gamma]\theta \wedge (\varsigma_x, v_x) \in \mathcal{V}_k[\theta(t)]\} \\
\mathcal{H}[\sigma] &= \star_{(l,f,m) \in \sigma} [l \mapsto_f m] \\
\varsigma &\equiv \mathcal{H}[\theta(\sigma)] \\
{}_k[\Theta; \Delta; \Gamma \vdash e : t] &= \forall \theta, \delta, \gamma, \sigma. \Theta = \text{dom}(\theta) \wedge (\varsigma, \gamma) \in \mathcal{L}_k[\Gamma]\theta \wedge \delta \in \mathcal{I}_k[\Delta]\theta \Rightarrow \\
&\quad (\varsigma, \theta(\delta(\gamma(e)))) \in \mathcal{C}_k[\theta(t)]
\end{aligned}$$

800 **C** Lemmas

801 **C.1** $\forall \sigma_s, \sigma_r, e. \varsigma_s \star \varsigma_r \text{ defined} \Rightarrow \forall n. \langle \sigma_s, e \rangle \rightarrow^n = \langle \sigma_s + \sigma_r, e \rangle \rightarrow^n$

802 SUFFICES: By induction on n , consider only the cases $\langle \sigma_s, e \rangle \rightarrow \langle \sigma_f, e_f \rangle$ where $\sigma_s \neq \sigma_f$.

803
804 PROOF SKETCH: Only `OP_FREE`, `MATRIX`, `SHARE`, `UNSHARE_EQ`, `GEMM_MATCH`
805 change the heap: the rest are either parametric in the heap or step to an **err**.

806
807 PROVE: $\langle \sigma_s + \sigma_r, e \rangle \rightarrow \langle \sigma_f + \sigma_r, e_f \rangle$.

808
809 $\langle 1 \rangle 1$. CASE: `OP_FREE`, $\sigma_s \equiv \sigma' + \{l \mapsto_1 m\}$, $\sigma_f = \sigma'$.

810 PROOF: Instantiate `OP_FREE` with $(\sigma' + \sigma_r) + \{l \mapsto_1 m\}$,
811 valid because $l \notin \text{dom}(\varsigma_r)$ by $\varsigma' \star [l \mapsto_1 m] \star \varsigma_r$ defined (assumption).

812 $\langle 1 \rangle 2$. CASE: `OP_MATRIX`

813 PROOF: Rule has no requirements on σ_s so will also work with $\sigma_s + \sigma_r$.

814 $\langle 1 \rangle 3$. CASE: `OP_SHARE`, $\sigma_s \equiv \sigma' + \{l \mapsto_f m\}$, $\sigma_f = \sigma' + \{l \mapsto_{\frac{1}{2}.f} m\} + \{l \mapsto_{\frac{1}{2}.f} m\}$.

815 PROOF: Union-ing σ_r does not remove $l \mapsto_f m$, so that can be split out of $\sigma_s + \sigma_r$ as
816 before.

817 $\langle 1 \rangle 4$. CASE: `OP_UNSHARE_EQ`, $\sigma_s \equiv \sigma' + \{l \mapsto_{\frac{1}{2}.f} m\} + \{l \mapsto_{\frac{1}{2}.f} m\}$, $\sigma_f = \sigma' + \{l \mapsto_f m\}$.

818 $\langle 2 \rangle 1$. Union-ing σ_r does not remove $l \mapsto_{\frac{1}{2}.f} m$, so that can still be split out of $\sigma_s + \sigma_r$.

819 $\langle 2 \rangle 2$. There may also be other valid splits introduced by σ_r .

820 $\langle 2 \rangle 3$. However, by assumption of $\varsigma_s \star \varsigma_r$ defined, any splitting of $\sigma_s + \sigma_r$ will satisfy
821 $f \leq 1$.

822 $\langle 1 \rangle 5$. CASE: `OP_GEMM_MATCH`

823 $\langle 2 \rangle 1$. By assumption of $\varsigma_s \star \varsigma_r$ defined, either l_1 (or l_2 , or both) are not in σ_r , or they
824 are and the matrix values they point to are the same.

825 $\langle 2 \rangle 2$. The permissions (of l_1 and/or l_2) may differ, but `OP_GEMM_MATCH` universally
826 quantifies over them and leaves them unchanged, so they are irrelevant.

827 $\langle 2 \rangle 3$. Only the pointed to matrix value at l_3 changes.

828 $\langle 2 \rangle 4$. SUFFICES: $l_3 \notin \pi_1[\sigma_r]$.

829 $\langle 2 \rangle 5$. By assumption of $\varsigma_s \star \varsigma_r$ defined, $l_3 \notin \text{dom}(\varsigma_r)$.

830 $\langle 2 \rangle 6$. Hence $l_3 \notin \pi_1[\sigma_r]$.

831 **C.2** $\forall k, t. \mathcal{V}_k[t] \subseteq \mathcal{C}_k[t]$

832 Follows from definition of $\mathcal{C}_k[t]$, \rightarrow^j ($\forall n. \langle \sigma, v \rangle \rightarrow^n \langle \sigma, v \rangle$) for arbitrary $j \leq k$ and C.1.

833 **C.3** $\forall \theta, \delta, \gamma, v. \theta(\delta(\gamma(v)))$ is a value.

834 θ is irrelevant because it only maps fractional permission variables to fractional permissions.

835 By construction, δ and γ only map variables to values, and values are closed under substitution.

836 **C.4** $\forall k, \sigma, \sigma', e, e', t. (\varsigma', e') \in \mathcal{C}_k[t] \wedge \langle \sigma, e \rangle \rightarrow \langle \sigma', e' \rangle \Rightarrow (\varsigma, e) \in \mathcal{C}_{k+1}[t]$

837 In the lemma, and for the rest of its proof, $\varsigma = \mathcal{H}[\sigma]$.

838 ASSUME: arbitrary $j < k + 1$, and σ_r such that $\varsigma \star \varsigma_r$ defined.

839

840 $\langle 1 \rangle 1$. CASE: $j = 0$. Clearly $\sigma_f = \sigma_s + \sigma_r$ and $e' = e$.

841 Remains to show that if e is a value then $(\varsigma_s \star \varsigma_r, e) \in \mathcal{V}_k[t]$.

842 This is true vacuously, because by assumption, e is not a value.

843 $\langle 1 \rangle 2$. CASE: $j \geq 1$. We have $\langle \sigma, e \rangle \rightarrow^j = \langle \sigma', e' \rangle \rightarrow^{j-1}$.

844 Instantiate $(\varsigma', e') \in \mathcal{C}_k[t]$, with $j - 1 < k$ and σ_r to conclude the required conditions.

845 **C.5** $j \leq k \Rightarrow _k[\cdot] \subseteq _j[\cdot]$

846 For the rest of this proof, $\varsigma = \mathcal{H}[\sigma]$.

847 Lemma C.4 is the inductive step for this lemma for the $\mathcal{C}[\cdot]$ case.

848 Need to prove for $\mathcal{V}[\cdot]$, by induction on t and then index.

849 SUFFICES: Consider only $t \multimap t'$ case, rest use k directly on structure of type.

850 ASSUME: Arbitrary $j \leq k$ and $(\varsigma_{v'}, v') \in \mathcal{V}_k[t \multimap t']$.

851 PROVE: $(\varsigma_{v'}, v') \in \mathcal{V}_j[t \multimap t']$.

852

853 $\langle 1 \rangle 1$. v' is of the correct syntactic form (lambda or fixpoint) by assumption.

854 $\langle 1 \rangle 2$. ASSUME: arbitrary $j' \leq j$ and $(\varsigma_v, v) \in \mathcal{V}_{j'}[t]$ such that $\varsigma_{v'} \star \varsigma_v$ is defined.

855 $\langle 1 \rangle 3$. SUFFICES: to show $(\varsigma_{v'} \star \varsigma_v, v'v) \in \mathcal{C}_{j'}[t']$.

856 $\langle 1 \rangle 4$. This is true by instantiating $(\varsigma_{v'}, v') \in \mathcal{V}_k[t \multimap t']$ with $j' \leq k$ and $(\varsigma_v, v) \in \mathcal{V}_{j'}[t]$.

857 **C.6** $\forall \Delta, \Gamma, t, k, \theta, \delta, \gamma. \delta \in \mathcal{I}_k[\Delta]\theta \wedge \gamma \in \pi_2[\mathcal{L}_k[\Gamma]\theta] \Rightarrow \text{dom}(\Delta) = \text{dom}(\delta)$
 858 **and** $\text{dom}(\Gamma) = \text{dom}(\gamma)$

859 PROOF: By induction on Δ and Γ .

860 **C.7** $\forall k, \Gamma, \Gamma', \theta, \sigma_+, \gamma_+. (\varsigma_+, \gamma) \in \mathcal{L}_k[\Gamma, \Gamma']\theta \wedge \Gamma, \Gamma' \text{ disjoint} \Rightarrow$
 861 $\exists \sigma, \gamma, \sigma', \gamma'. \sigma_+ = \sigma + \sigma' \wedge \gamma, \gamma' \text{ disjoint} \wedge \gamma_+ = \gamma \cup \gamma' \wedge (\varsigma, \gamma) \in$
 862 $\mathcal{L}_k[\Gamma] \wedge (\varsigma', \gamma') \in \mathcal{L}_k[\Gamma']$

863 PROOF: By induction on Γ' .

23:30 NumLin: Linear Types for Linear Algebra

864 **C.8** $\forall e, \sigma, e', \sigma', \theta. \langle \sigma, e \rangle \rightarrow \langle \sigma', e' \rangle \Rightarrow \langle \theta(\sigma), \theta(e) \rangle \rightarrow \langle \theta(\sigma'), \theta(e') \rangle$

865 PROOF: By induction on \rightarrow .

866 $\langle 1 \rangle 1$. ASSUME: Arbitrary $e, \sigma, e', \sigma', \theta$ such that $\langle \sigma, e \rangle \rightarrow \langle \sigma', e' \rangle$.

867 $\langle 1 \rangle 2$. SUFFICES: To consider only the following rules which mention fractional permission
868 variables.

869 OP_FRAC_PERM, OP_SHARE, OP_UNSHARE_(N)EQ and OP_GEMM_(MIS)MATCH.

870 $\langle 1 \rangle 3$. CASE: OP_FRAC_PERM.

871 Because substitution avoids capture,

872 $\langle \theta(\sigma), \theta(\text{fun } fc \rightarrow v) [f] \rangle \rightarrow \langle \theta(\sigma' [fc/f]), \theta(v [fc/f]) \rangle$.

873 $\langle 1 \rangle 4$. The rest of the cases are parametric in their use of fractional permission variables and
874 so will take the same step after any substitution.

875 $\langle 1 \rangle 5$. COROLLARY: If $\langle \sigma [fc/f_1], e [fc/f_1] \rangle \rightarrow^n \langle \sigma_2, e'_2 \rangle$ and $\langle \sigma [fc/f_2], e [fc/f_2] \rangle \rightarrow^n \langle \sigma_2, e'_2 \rangle$,
876 then $\exists \sigma, e'. \sigma_1 = \sigma [fc/f_1] \wedge \sigma_2 = \sigma [fc/f_2] \wedge e'_1 = e' [fc/f_1] \wedge e'_2 = e' [fc/f_2]$.

877 **D Soundness**

878 $\forall \Theta, \Delta, \Gamma, e, t. \Theta; \Delta; \Gamma \vdash e : t \Rightarrow \forall k. {}_k \llbracket \Theta; \Delta; \Gamma \vdash e : t \rrbracket$

879 PROOF SKETCH: Induction over the typing judgements.

880

881 ASSUME: 1. Arbitrary $\Theta, \Delta, \Gamma, e, t$ such that $\Theta; \Delta; \Gamma \vdash e : t$.

882 2. Arbitrary $k, \theta, \delta, \gamma, \sigma$ such that:

883 a. $\Theta = \text{dom}(\theta)$

884 b. $\delta \in \mathcal{I}_k \llbracket \Delta \rrbracket \theta$.

885 c. $(\varsigma, \gamma) \in \mathcal{L}_k \llbracket \Gamma \rrbracket \theta$

886 3. W.l.o.g., all variables are distinct, hence Θ , $\text{dom}(\Delta)$ and $\text{dom}(\Gamma)$ are disjoint so
887 order of θ , δ and γ (as substitutions defined recursively over expressions) is
888 irrelevant.

889

890 PROVE: $(\varsigma, \theta(\delta(\gamma(e)))) \in \mathcal{C}_k \llbracket \theta(t) \rrbracket$.

891 ASSUME: Arbitrary $j < k$ and σ_r , such that $\varsigma \star \varsigma_r$ defined.

892 SUFFICES: $\langle \sigma + \sigma_r, e \rangle \rightarrow^j \text{err} \vee \exists \sigma_f, e_f. \langle \sigma + \sigma_r, e \rangle \rightarrow^j \langle \sigma_f + \sigma_r, e_f \rangle$

893 $\wedge (e_f \text{ is a value} \Rightarrow (\varsigma_f \star \varsigma_r, e_f) \in \mathcal{V}_{k-j} \llbracket t \rrbracket)$.

894 SUFFICES: By C.1, to show $\langle \sigma, e \rangle \rightarrow^j \text{err} \vee \exists \sigma_f, e_f. \langle \sigma, e \rangle \rightarrow^j \langle \sigma_f, e_f \rangle$

895 $\wedge (e_f \text{ is a value} \Rightarrow (\varsigma_f, e_f) \in \mathcal{V}_{k-j} \llbracket t \rrbracket)$

896

897 $\langle 1 \rangle 1$. CASE: TY_LET.

898 $\langle 2 \rangle 1$. By induction,

899 1. $\forall k. {}_k \llbracket \Theta; \Delta; \Gamma \vdash e : t \rrbracket$

900 2. $\forall k. {}_k \llbracket \Theta; \Delta; \Gamma', x : t \vdash e' : t' \rrbracket$.

- 901 $\langle 2 \rangle 2$. By 2c, 3 and C.7, we know there exists the following (for all k):
 902 1. $(\varsigma_e, \gamma_e) \in \mathcal{L}_k[\Gamma]$
 903 2. $\gamma = \gamma_e \cup \gamma_{e'}$
 904 3. $\sigma = \sigma_e + \sigma_{e'}$.
- 905 $\langle 2 \rangle 3$. So, using $k, \theta, \delta, \gamma_e, \sigma_e$, we have $(\varsigma_e, \theta(\delta(\gamma_e(e)))) \in \mathcal{C}_k[\theta(t)]$.
- 906 $\langle 2 \rangle 4$. By $\langle 2 \rangle 2$ ($\gamma = \gamma_e \cup \gamma_{e'}$), have $(\varsigma_e, \theta(\delta(\gamma(e)))) \in \mathcal{C}_k[\theta(t)]$.
- 907 $\langle 2 \rangle 5$. By definition of $\mathcal{C}_k[\cdot]$ and $\langle 2 \rangle 2$, we instantiate with j and $\sigma_r = \sigma_{e'}$ to conclude
 908 that
 909 $\langle \theta(\sigma), \theta(\delta(\gamma(e))) \rangle$ either takes j steps to **err** or another heap-and-expression
 910 $\langle \sigma_f, e_f \rangle$.
- 911 $\langle 2 \rangle 6$. CASE: j steps to **err**
 912 By `OP_CONTEXT_ERR`, the whole expression reduces to **err** in $j < k$ steps.
- 913 $\langle 2 \rangle 7$. CASE: j steps to another heap-and-expression.
 914 If it is not a value, then `OP_CONTEXT` runs j times and we are done.
- 915 $\langle 2 \rangle 8$. If it is, then $\exists i \leq j. (\varsigma_f, v_1) \in \mathcal{V}_{k-i}[\theta(t_1)] \subseteq \mathcal{V}_{k-j}[\theta(t_1)]$ by C.3 and C.5.
 916 So, `OP_CONTEXT` runs i times, and then we have the following.
 917 SUFFICES: $(\varsigma_f \star \varsigma_{e'}, \mathbf{let} \ x = v \ \mathbf{in} \ \theta(\delta(\gamma(e')))) \in \mathcal{C}_{k-i}[\theta(t')]$ by C.4 i times.
 918 SUFFICES: $(\varsigma_f \star \varsigma_{e'}, \theta(\delta(\gamma(e')))[x/v]) \in \mathcal{C}_{k-i-1}[\theta(t')]$ by C.4.
- 919 $\langle 2 \rangle 9$. By C.5, $(\varsigma_{e'}, \gamma_{e'}[x \mapsto v]) \in \mathcal{L}_k[\Gamma', x : t]\theta \subseteq \mathcal{L}_{k-i-1}[\Gamma', x : t]\theta$.
- 920 $\langle 2 \rangle 10$. Instantiate 2 of step $\langle 2 \rangle 1$ with $k - i - 1, \theta, \delta, \gamma_{e'}[x \mapsto v], \sigma_{e'}$ to conclude
 921 $(\varsigma_{e'}, \theta(\delta(\gamma_{e'}[x \mapsto v](e')))) \in \mathcal{C}_{k-i-1}[\theta(t')]$.
- 922 $\langle 2 \rangle 11$. By 3, we have $\theta(\delta(\gamma(e')))[x/v] = \theta(\delta(\gamma_{e'}[x \mapsto v](e')))$ and
 923 by C.1 we conclude $(\varsigma_f \star \varsigma_{e'}, \theta(\delta(\gamma(e')))[x/v]) \in \mathcal{C}_{k-i-1}[\theta(t')]$.
- 924 $\langle 1 \rangle 2$. CASE: `TY_PAIR_ELIM`.
 925 PROOF SKETCH: Similar to `TY_LET`, but with the following key differences.
- 926 $\langle 2 \rangle 1$. When $(\varsigma_f, v) \in \mathcal{V}_{k-i}[\theta(t_1) \otimes \theta(t_2)]$, we have $v = (v_1, v_2)$.
- 927 $\langle 2 \rangle 2$. SUFFICES: $(\varsigma_{e'}, \theta(\delta(\gamma(e')))) \in \mathcal{C}_{k-i-1}[\theta(t')]$ by C.4 $i + 1$ times.
- 928 $\langle 2 \rangle 3$. By C.5, $(\varsigma_{e'}, \gamma_{e'}[a \mapsto v_1, b \mapsto v_2]) \in \mathcal{L}_k[\Gamma', a : t_1, b : t_2]\theta \subseteq \mathcal{L}_{k-i-1}[\Gamma', a : t_1, b : t_2]\theta$.
 929 $t_2]\theta$.
- 930 $\langle 2 \rangle 4$. Instantiate $_{k-i-1}[\Theta; \Delta; \Gamma', a : t_1, b : t_2 \vdash e' : t']$ with $\theta, \delta, \gamma_{e'}[a \mapsto v_1, b \mapsto v_2], \sigma_{e'}$.
- 931 $\langle 2 \rangle 5$. By 3 (for $\gamma = \gamma_e \cup \gamma_{e'}$ and a, b), conclude $(\varsigma_{e'}, \theta(\delta(\gamma(e'[a/v_1][b/v_2]))) \in \mathcal{C}_{k-i-1}[\theta(t')]$.
- 932 $\langle 1 \rangle 3$. CASE: `TY_BANG_ELIM`.
 933 PROOF SKETCH: Similar to `TY_LET`, but with the following key differences.
- 934 $\langle 2 \rangle 1$. When $(\varsigma_f, v) \in \mathcal{V}_{k-i}[\theta(!t)]$, since $\mathcal{V}_{k-i}[\theta(!t)] = \mathcal{V}_{k-i}[\theta(t)]$,
 935 we have $\varsigma_f = \emptyset$ and $v = \mathbf{Many} \ v'$ for some $(\emptyset, v') \in \mathcal{V}_{k-i}[\theta(t)]$.

- 936 $\langle 2 \rangle 2$. SUFFICES: $(\varsigma_{e'}, \mathbf{let\ Many\ } x = \mathbf{Many\ } v' \mathbf{ in\ } \theta(\delta(\gamma(e')))) \in \mathcal{C}_{k-i}[\![\theta(t)]\!]$.
- 937 $\langle 2 \rangle 3$. SUFFICES: $(\varsigma_{e'}, \theta(\delta(\gamma(e')))[x/v]) \in \mathcal{C}_{k-i-1}[\![\theta(t)]\!]$ by C.4 $i + 1$ times.
- 938 $\langle 2 \rangle 4$. Instantiate ${}_{k-i-1}[\![\Theta; \Delta, x : t, \Gamma' \vdash e' : t']\!]$ with $\theta, \delta_{e'} = \delta[x \mapsto v'], \gamma_{e'}, \sigma_{e'}$.
- 939 $\langle 2 \rangle 5$. By 3, $(\varsigma_{e'}, \theta(\delta(\gamma(e')))[x/v]) \in \mathcal{C}_{k-i-1}[\![\theta(t)]\!]$.
- 940 $\langle 1 \rangle 4$. CASE: TY_UNIT_ELIM.
- 941 PROOF SKETCH: Similar to TY_LET, but with the following key differences.
- 942 $\langle 2 \rangle 1$. When $(\varsigma_f, v) \in \mathcal{V}_{k-i}[\![\mathbf{unit}]\!]$, we have $\varsigma_f = \emptyset$ and $v = ()$.
- 943 $\langle 2 \rangle 2$. SUFFICES: $(\varsigma_{e'}, \theta(\delta(\gamma(e')))) \in \mathcal{C}_{k-i-1}[\![\theta(t')]\!]$ by C.4 $i + 1$ times.
- 944 $\langle 2 \rangle 3$. By C.5, $(\varsigma_{e'}, \gamma_{e'}) \in \mathcal{L}_k[\![\Gamma']\!]\theta \subseteq \mathcal{L}_{k-i-1}[\![\Gamma']\!]\theta$.
- 945 $\langle 2 \rangle 4$. Instantiate ${}_{k-i-1}[\![\Theta; \Delta; \Gamma' \vdash e' : t']\!]$ with $\theta, \delta, \gamma_{e'}, \sigma_{e'}$.
- 946 $\langle 2 \rangle 5$. By 3 $(\varsigma_{e'}, \theta(\delta(\gamma(e')))) \in \mathcal{C}_{k-i-1}[\![\theta(t')]\!]$.
- 947 $\langle 1 \rangle 5$. CASE: TY_BOOL_ELIM.
- 948 PROOF SKETCH: Similar to TY_UNIT_ELIM but with $\mathbf{OP_IF_}\{\mathbf{TRUE}, \mathbf{FALSE}\}$, $\varsigma_f = \emptyset$
- 949 and $v = \mathbf{Many\ true}$ or $v = \mathbf{Many\ false}$.
- 950 $\langle 1 \rangle 6$. CASE: TY_BANG_INTRO.
- 951 $\langle 2 \rangle 1$. We have, $e = v$ for some value $v \neq l$, $\Gamma = \emptyset$ and so
- 952 $\forall k. {}_k[\![\Theta; \Delta; \cdot \vdash v : t]\!]$ by induction.
- 953 $\langle 2 \rangle 2$. SUFFICES: $(\emptyset, \mathbf{Many\ } \theta(\delta(v))) \in \mathcal{C}_k[\![!\theta(t)]\!]$ by 2c ($\varsigma = \emptyset, \gamma = []$).
- 954 $\langle 2 \rangle 3$. Instantiate ${}_k[\![\Theta; \Delta; \cdot \vdash v : t]\!]$ with $\theta, \delta, \gamma = [], \sigma = \emptyset$ to obtain $(\emptyset, \theta(\delta(v))) \in$
- 955 $\mathcal{C}_k[\![\theta(t)]\!]$.
- 956 $\langle 2 \rangle 4$. Instantiate $(\emptyset, \theta(\delta(v))) \in \mathcal{C}_k[\![\theta(t)]\!]$ with $j = 0$, $\sigma_r = \emptyset$ and C.3 ($\theta(\delta(v))$ is a value),
- 957 to conclude $(\emptyset, \theta(\delta(v))) \in \mathcal{V}_k[\![\theta(t)]\!]$.
- 958 $\langle 2 \rangle 5$. By definition of $\mathcal{V}_k[\![!\theta(t)]\!]$, C.3 and C.2 we have $(\emptyset, \mathbf{Many\ } \theta(\delta(v))) \in \mathcal{C}_k[\![!\theta(t)]\!]$.
- 959 $\langle 1 \rangle 7$. CASE: TY_PAIR_INTRO.
- 960 $\langle 2 \rangle 1$. By 2c, 3 and C.7, we know there exists the following (for all k):
- 961 1. $(\varsigma_1, \gamma_1) \in \mathcal{L}_k[\![\Gamma_1]\!]$
- 962 2. $(\varsigma_2, \gamma_2) \in \mathcal{L}_k[\![\Gamma_2]\!]$
- 963 3. $\gamma = \gamma_1 \cup \gamma_2$
- 964 4. $\sigma = \sigma_1 + \sigma_2$.
- 965 $\langle 2 \rangle 2$. By induction,
- 966 1. $\forall k. {}_k[\![\Theta; \Delta; \Gamma_1 \vdash e_1 : t_1]\!]$
- 967 2. $\forall k. {}_k[\![\Theta; \Delta; \Gamma_2 \vdash e_2 : t_2]\!]$.
- 968 $\langle 2 \rangle 3$. Instantiate the first with $k, \theta, \delta, \gamma_1, \sigma_1$.

- 969 $\langle 2 \rangle 4$. By that and $\langle 2 \rangle 1$, $(\varsigma_1, \theta(\delta(\gamma_1(e_1)))) = (\varsigma_1, \theta(\delta(\gamma(e_1)))) \in \mathcal{C}_k[\![\theta(t)]\!]$.
- 970 $\langle 2 \rangle 5$. So, $\langle \theta(\sigma_1 + \sigma_2), \theta(\delta(\gamma_1(e_1))) \rangle$ either takes j steps to **err** or a heap-and-expression
 971 $\langle \sigma_{1f}, e_{1f} \rangle$.
- 972 $\langle 2 \rangle 6$. CASE: j steps to **err**
 973 By `OP_CONTEXT_ERR`, the whole expression reduces to **err** in $j < k$ steps.
- 974 $\langle 2 \rangle 7$. CASE: j steps to another heap-and-expression.
 975 If it is not a value, then `OP_CONTEXT` runs j times and we are done.
- 976 $\langle 2 \rangle 8$. If it is, then $\exists i_1 \leq j. (\varsigma_{1f}, v_1) \in \mathcal{V}_{k-i_1}[\![\theta(t_1)]\!]$ by C.3 and C.5.
 977 So, `OP_CONTEXT` runs i_1 times, and then we have the following.
 978 SUFFICES: By C.4, $(\varsigma_{1f} \star \varsigma_2, (v_1, e_2)) \in \mathcal{C}_{k-i_1}[\![\theta(t_1) \otimes t_2]\!]$.
- 979 $\langle 2 \rangle 9$. Instantiate the second IH with $k, \theta, \delta, \gamma_2, \sigma_2$.
- 980 $\langle 2 \rangle 10$. So, $\langle \theta(\sigma_{1f} + \sigma_2), \theta(\delta(\gamma_2(e_2))) \rangle$ either takes j steps to **err** or a heap-and-expression
 981 $\langle \sigma_{2f}, e_{2f} \rangle$.
- 982 $\langle 2 \rangle 11$. CASE: j steps to **err**
 983 By `OP_CONTEXT_ERR`, the whole expression reduces to **err** in $j < k$ steps.
- 984 $\langle 2 \rangle 12$. CASE: j steps to another heap-and-expression.
 985 If it is not a value, then `OP_CONTEXT` runs j times and we are done.
- 986 $\langle 2 \rangle 13$. If it is, then $\exists i_2 \leq j. (\varsigma_{2f}, v_2) \in \mathcal{V}_{k-i_2}[\![\theta(t_2)]\!]$ by C.3 and C.5.
 987 So, `OP_CONTEXT` runs i_2 times, and then we have the following.
 988 SUFFICES: By C.4, $(\varsigma_{1f} \star \varsigma_{2f}, (v_1, v_2)) \in \mathcal{V}_{k-i_1-i_2}[\![\theta(t_1) \otimes \theta(t_2)]\!]$.
- 989 $\langle 2 \rangle 14$. By C.5 and $k - i_1 - i_2 \leq k - i_1, k - i_2$, have
 990 $(\varsigma_{1f}, v_1) \in \mathcal{V}_{k-i_1}[\![\theta(t_1)]\!]$ and
 991 $(\varsigma_{2f}, v_2) \in \mathcal{V}_{k-i_2}[\![\theta(t_2)]\!]$ as needed.
- 992 $\langle 1 \rangle 8$. CASE: `TY_LAMBDA`.
 993 SUFFICES: By C.2, to show $(\varsigma, \theta(\delta(\gamma(\mathbf{fun} x : t \rightarrow e)))) \in \mathcal{V}_k[\![\theta(t \rightarrow t')]\!]$.
 994 ASSUME: Arbitrary $j \leq k$, $(\varsigma_v, v) \in \mathcal{V}_j[\![\theta(t)]\!]$ such that $\varsigma \star \varsigma_v$ is defined.
 995 SUFFICES: $(\varsigma \star \varsigma_v, \theta(\delta(\gamma(\mathbf{fun} x : t \rightarrow e))) v) \in \mathcal{C}_j[\![\theta(t')]\!]$.
 996 SUFFICES: $(\varsigma \star \varsigma_v, \theta(\delta(\gamma(e)))[x/v]) \in \mathcal{C}_{j-1}[\![\theta(t')]\!]$ by C.4.
- 997 $\langle 2 \rangle 1$. By induction, $\forall k. {}_k\llbracket \Theta; \Delta; \Gamma, x : t \vdash e \rrbracket$.
- 998 $\langle 2 \rangle 2$. Instantiate it $j - 1, \theta, \delta, \gamma[x \mapsto v], \sigma + \sigma_v$.
- 999 $\langle 2 \rangle 3$. Hence, $(\varsigma \star \varsigma_v, \theta(\delta(\gamma[x \mapsto v](e)))) \in \mathcal{C}_{j-1}[\![\theta(t')]\!]$.
- 1000 $\langle 2 \rangle 4$. By 3, $\theta(\delta(\gamma[x \mapsto v](e))) = \theta(\delta(\gamma(e)))[x/v]$, we are done.
- 1001 $\langle 1 \rangle 9$. CASE: `TY_APP`.
- 1002 $\langle 2 \rangle 1$. By 2c, 3 and C.7, we know there exists the following (for all k):
 1003 1. $(\varsigma_e, \gamma_e) \in \mathcal{L}_k[\![\Gamma_e]\!]$

- 1004 2. $(\varsigma_{e'}, \gamma_{e'}) \in \mathcal{L}_k[\Gamma_{e'}]$
 1005 3. $\gamma = \gamma_e \cup \gamma_{e'}$
 1006 4. $\sigma = \sigma_e + \sigma_{e'}$.
- 1007 $\langle 2 \rangle 2$. By induction,
 1008 1. $\forall k. {}_k[\Theta; \Delta; \Gamma \vdash e : t' \multimap t]$
 1009 2. $\forall k. {}_k[\Theta; \Delta; \Gamma' \vdash e' : t']$.
- 1010 $\langle 2 \rangle 3$. Instantiate the first with $k, \theta, \delta, \gamma_e, \sigma_e$ to conclude $(\varsigma_e, \theta(\delta(\gamma_e(e)))) \in \mathcal{C}_k[\theta(t') \multimap \theta(t)]$.
 1011
- 1012 $\langle 2 \rangle 4$. Instantiate *this* with j and $\sigma_{e'}$ and use $\langle 2 \rangle 1$ to conclude $\langle \theta(\sigma_e + \sigma_{e'}), \theta(\delta(\gamma_e(e))) \rangle$
 1013 either takes j steps to **err** or a heap-and-expression $\langle \sigma_f + \sigma_{e'}, e_f \rangle$.
- 1014 $\langle 2 \rangle 5$. CASE: j steps to **err**
 1015 By OP_CONTEXT_ERR, the whole expression reduces to **err** in $j < k$ steps.
- 1016 $\langle 2 \rangle 6$. CASE: j steps to another heap-and-expression.
 1017 If it is not a value, then OP_CONTEXT runs j times and we are done.
- 1018 $\langle 2 \rangle 7$. If it is, then $\exists i_e \leq j. (\varsigma_f, e_f) \in \mathcal{V}_{k-i_e}[\theta(t') \multimap \theta(t)] \subseteq \mathcal{V}_{k-j}[\dots]$ by C.3 and C.5.
 1019 So, OP_CONTEXT runs i_e times, and then we have the following.
 1020 SUFFICES: By C.4 i_e times, $(\varsigma_f \star \varsigma_{e'}, e_f e') \in \mathcal{C}_{k-i_e}[\theta(t')]$.
- 1021 $\langle 2 \rangle 8$. By C.5, $(\varsigma_{e'}, \gamma_{e'}) \in \mathcal{L}_k[\Gamma'] \theta \subseteq \mathcal{L}_{k-i_e}[\Gamma'] \theta$.
- 1022 $\langle 2 \rangle 9$. So, instantiate the second IH with $k - i_e, \theta, \delta, \gamma_{e'}, \sigma_{e'}$ to conclude
 1023 $(\varsigma_{e'}, \theta(\delta(\gamma_{e'}(e')))) \in \mathcal{C}_{k-i_e}[\theta(t')]$.
- 1024 $\langle 2 \rangle 10$. Instantiate *this* with $j - i_e$ and σ_f to conclude $\langle \theta(\sigma_f + \sigma_{e'}), \theta(\delta(\gamma_{e'}(e')) \rangle$
 1025 either takes $j - i_e$ steps to **err** or $\langle \sigma_f + \sigma'_f, e'_f \rangle$.
- 1026 $\langle 2 \rangle 11$. CASE: $j - i_e$ steps to **err**
 1027 By OP_CONTEXT_ERR, the whole expression reduces to **err** in $j - i_e < k - i_e$
 1028 steps.
- 1029 $\langle 2 \rangle 12$. CASE: $j - i_e$ steps to another heap-and-expression.
 1030 If it is not a value, then OP_CONTEXT runs $j - i_e$ times and we are done.
- 1031 $\langle 2 \rangle 13$. If it is, then $\exists i_{e'} \leq j - i_e. (\varsigma'_f, v_{e'}) \in \mathcal{V}_{k-i_e-i_{e'}}[\theta(t')]$ by C.3.
 1032 So, OP_CONTEXT runs $i_{e'}$ times, and then we have the following.
 1033 SUFFICES: By C.4 $i_{e'}$ times, $(\varsigma_f \star \varsigma'_f, e_f e'_f) \in \mathcal{C}_{k-i_e-i_{e'}}[\theta(t')]$.
- 1034 $\langle 2 \rangle 14$. Instantiate $(\varsigma_f, e_f) \in \mathcal{V}_{k-i_e}[\theta(t') \multimap \theta(t)]$ with $k - i_e - i_{e'} \leq k - i_e$ and
 1035 $(\varsigma_{v'}, v_{e'}) \in \mathcal{V}_{k-i_e-i_{e'}}[\theta(t')]$, to conclude $(\varsigma_f \star \varsigma'_f, e_f e'_f) \in \mathcal{C}_{k-i_e-i_{e'}}[\theta(t)]$ as
 1036 needed.
- 1037 $\langle 1 \rangle 10$. CASE: TY_GEN.
- 1038 $\langle 2 \rangle 1$. By induction, $\forall k. {}_k[\Theta, fc; \Delta; \Gamma \vdash e : t]$.
- 1039 $\langle 2 \rangle 2$. LET: f be arbitrary; $\theta' \equiv \theta[fc \mapsto f]$.

1040 Instantiate induction hypothesis with $k - 1, \theta', \delta, \gamma, \sigma$,
 1041 to conclude $(\varsigma, \theta'(\gamma(\delta(e)))) \in \mathcal{C}_{k-1}[\![\theta'(t)]\!]$ (for all f , by C.8).

1042 $\langle 2 \rangle 3$. Instantiate *this* with j and \emptyset to conclude $\langle \theta(\sigma), \theta'(\gamma(\delta(e))) \rangle$
 1043 either takes j steps to **err** or a heap-and-expression $\langle \sigma', e' \rangle$ (for all f , by C.8).

1044 $\langle 2 \rangle 4$. CASE: j steps to **err**.
 1045 By OP_CONTEXT_ERR, whole expression reduces to **err** in $j < k - 1$ steps
 1046 (for $f = fc$).

1047 $\langle 2 \rangle 5$. CASE: j steps to another heap-and-expression.
 1048 If it is not a value, then for $f = fc$, OP_CONTEXT runs j times and we are
 1049 done.

1050 $\langle 2 \rangle 6$. If it is, then $\exists i_e \leq j. (\varsigma', e') \in \mathcal{V}_{k-1-i_e}[\![\theta'(t)]\!] \subseteq \mathcal{V}_{k-1-j}[\![\dots]\!]$
 1051 by C.3 and C.5 (for all f , by C.8).

1052 $\langle 2 \rangle 7$. So, OP_CONTEXT runs i_e times, and then we have the following.
 1053 SUFFICES: By C.4 i_e times, $(\varsigma', \mathbf{fun} fc \rightarrow e') \in \mathcal{V}_{k-i_e}[\![\theta(\forall fc. t)]\!]$ (for $f = fc$).

1054 $\langle 2 \rangle 8$. ASSUME: Arbitrary f' .
 1055 SUFFICES: $(\varsigma', e'[fc/f']) \in \mathcal{V}_{k-1-i_e}[\![\theta(t)[fc/f']]\!]$ (for $f = fc$).

1056 $\langle 2 \rangle 9$. This is true by instantiating $\langle 2 \rangle 6$ with $f = f'$.

1057 $\langle 1 \rangle 11$. CASE: TY_SPC.

1058 $\langle 2 \rangle 1$. By induction, $\forall k. {}_k[\![\Theta; \Delta; \Gamma \vdash e : \forall fc. t]\!]$.

1059 $\langle 2 \rangle 2$. Instantiate with $k, \theta, \delta, \gamma, \sigma$ to conclude $(\varsigma, \theta(\delta(\gamma(e)))) \in \mathcal{C}_k[\![\theta(\forall fc. t)]\!]$.

1060 $\langle 2 \rangle 3$. Instantiate *this* with j and \emptyset and to conclude $\langle \theta(\sigma), \theta(\delta(\gamma(e))) \rangle$
 1061 either takes j steps to **err** or a heap-and-expression $\langle \sigma_f, e_f \rangle$.

1062 $\langle 2 \rangle 4$. CASE: j steps to **err**.
 1063 By OP_CONTEXT_ERR, the whole expression reduces to **err** in $j < k$ steps.

1064 $\langle 2 \rangle 5$. CASE: j steps to another heap-and-expression.
 1065 If it is not a value, then OP_CONTEXT runs j times and we are done.

1066 $\langle 2 \rangle 6$. If it is, then $\exists i_e \leq j. (\varsigma_f, e_f) \in \mathcal{V}_{k-i_e}[\![\theta(\forall fc. t)]\!] \subseteq \mathcal{V}_{k-j}[\![\dots]\!]$ by C.3 and C.5.
 1067 So $e_f \equiv \mathbf{fun} fc \rightarrow v$ for some v .

1068 $\langle 2 \rangle 7$. So, OP_CONTEXT runs i_e times, and then we have the following.
 1069 SUFFICES: By C.4 i_e times, $(\varsigma_f, (\mathbf{fun} fc \rightarrow v)[f]) \in \mathcal{C}_{k-i_e}[\![\theta(t[fc/f])]\!]$.
 1070 SUFFICES: By C.4 once more, $(\varsigma_f, v[fc/f]) \in \mathcal{C}_{k-i_e-1}[\![\theta(t[fc/f])]\!]$.

1071 $\langle 2 \rangle 8$. This is true by instantiating $\langle 2 \rangle 6$ with f and C.2.

1072 $\langle 1 \rangle 12$. CASE: TY_FIX.
 1073 SUFFICES: $(\emptyset, \theta(\delta(\mathbf{fix}(g, x : t, e : t')))) \in \mathcal{V}_k[\![\theta(t \multimap t')]\!]$, by C.2 ($\sigma = \{\}$, $\gamma = []$).
 1074 ASSUME: Arbitrary $j \leq k$, $(\varsigma_v, v) \in \mathcal{V}_j[\![\theta(t)]\!]$ ($\varsigma = \emptyset$, so $\varsigma \star \varsigma_v$ is defined).

23:36 NumLin: Linear Types for Linear Algebra

1075 LET: $\tilde{e} \equiv \theta(\delta(e))$.
 1076 SUFFICES: $(\varsigma_v, \mathbf{fix}(g, x : t, \tilde{e} : t') \ v) \in \mathcal{C}_j[\![\theta(t')]\!]$.
 1077 SUFFICES: $(\varsigma_v, \tilde{e} [x/v] [g/\mathbf{fix}(g, x : t, \tilde{e} : t')]) \in \mathcal{C}_{j-1}[\![\theta(t')]\!]$ by C.4.

1078 $\langle 2 \rangle 1$. By induction, $\forall k. {}_k[\![\Theta; \Delta, g : t \multimap t'; x : t \vdash e : t']]\!]$.
 1079 $\langle 2 \rangle 2$. Instantiate this with $j - 1, \delta[g \mapsto \mathbf{fix}(g, x : t, \tilde{e} : t')], \gamma = [x \mapsto v], \sigma_v$.
 1080 $\langle 2 \rangle 3$. We have $(\emptyset, \mathbf{fix}(g, x : t, \tilde{e} : t')) \in \mathcal{V}_{j-1}[\![\theta(t \multimap t')]\!]$.

1081 $\langle 3 \rangle 1$. Again by induction (over k), $(\emptyset, \mathbf{fix}(g, x : t, \tilde{e} : t')) \in \mathcal{C}_{j-1}[\![\theta(t \multimap t')]\!]$.
 1082 $\langle 3 \rangle 2$. Instantiate *this* with $j = 0$ and \emptyset and we are done.

1083 $\langle 2 \rangle 4$. We have $(\varsigma_v, v) \in \mathcal{V}_{j-1}[\![\theta(t)]\!]$ by assumption and C.5.
 1084 $\langle 2 \rangle 5$. So we conclude $(\varsigma_v, \theta(\delta'(\gamma(e)))) \in \mathcal{C}_{j-1}[\![\theta(t')]\!]$ as required.

1085 $\langle 1 \rangle 13$. CASE: `TY_VAR_LIN`.
 1086 PROVE: $(\varsigma, \theta(\delta(\gamma(x)))) \in \mathcal{C}_k[\![\theta(t)]\!]$.

1087 $\langle 2 \rangle 1$. $\Gamma = \{x : t\}$ by assumption of `TY_VAR_LIN`.
 1088 $\langle 2 \rangle 2$. SUFFICES: $(\varsigma, \gamma(x)) \in \mathcal{C}_k[\![\theta(t)]\!]$ by 3 (θ and δ irrelevant).
 1089 $\langle 2 \rangle 3$. By 2c, there exist $(\varsigma_x, v_x) \in \mathcal{V}_k[\![\theta(t)]\!]$, such that $\varsigma = \varsigma_x$ and $\gamma = [x \mapsto v_x]$.
 1090 $\langle 2 \rangle 4$. Hence, $(\varsigma_x, v_x) \in \mathcal{C}_k[\![\theta(t)]\!]$, by C.2.

1091 $\langle 1 \rangle 14$. CASE: `TY_VAR`.
 1092 PROVE: $(\varsigma, \theta(\delta(\gamma(x)))) \in \mathcal{C}_k[\![\theta(t)]\!]$.

1093 $\langle 2 \rangle 1$. $x : t \in \Delta$ and $\Gamma = \emptyset$ by assumption of `TY_VAR`.
 1094 $\langle 2 \rangle 2$. SUFFICES: $(\emptyset, \delta(x)) \in \mathcal{C}_k[\![\theta(t)]\!]$ by 3.
 1095 $\langle 2 \rangle 3$. By 2b, there exists v_x such that $(\emptyset, v_x) \in \mathcal{V}_k[\![\theta(t)]\!]$ (θ irrelevant and γ empty).
 1096 $\langle 2 \rangle 4$. Hence, $(\emptyset, v_x) \in \mathcal{C}_k[\![\theta(t)]\!]$, by C.2.

1097 $\langle 1 \rangle 15$. CASE: `TY_UNIT_INTRO`.
 1098 True by C.2 and definition of $\mathcal{V}_k[\![\mathbf{unit}]\!]$.

1099 $\langle 1 \rangle 16$. CASE: `TY_BOOL_TRUE`, `TY_BOOL_FALSE`, `TY_INT_INTRO`, `TY_ELT_INTRO`.
 1100 Similar to `TY_UNIT_INTRO`.

1101 D.1 Well-formed types

1102 $\boxed{\Theta \vdash f \text{ Perm}}$ Well-formed fractional permissions

1103 $\frac{fc \in \Theta}{\Theta \vdash fc \text{ Perm}}$ `WF_PERM_VAR`

1104	$\frac{}{\Theta \vdash 1 \text{ Perm}}$	WF_PERM_ZERO
1105	$\frac{\Theta \vdash f \text{ Perm}}{\Theta \vdash \frac{1}{2}f \text{ Perm}}$	WF_PERM_SUCC
1106	$\boxed{\Theta \vdash t \text{ Type}}$	Well-formed types
1107	$\frac{}{\Theta \vdash \text{unit Type}}$	WF_TYPE_UNIT
1108	$\frac{}{\Theta \vdash \text{bool Type}}$	WF_TYPE_BOOL
1109	$\frac{}{\Theta \vdash \text{int Type}}$	WF_TYPE_INT
1110	$\frac{}{\Theta \vdash \text{elt Type}}$	WF_TYPE_ELT
1111	$\frac{\Theta \vdash f \text{ Perm}}{\Theta \vdash f \text{ arr Type}}$	WF_TYPE_ARRAY
1112	$\frac{\Theta \vdash t \text{ Type}}{\Theta \vdash !t \text{ Type}}$	WF_TYPE_BANG
1113	$\frac{\Theta, fc \vdash t \text{ Type}}{\Theta \vdash \forall fc. t \text{ Type}}$	WF_TYPE_GEN
1114	$\frac{\Theta \vdash t \text{ Type} \quad \Theta \vdash t' \text{ Type}}{\Theta \vdash t \otimes t' \text{ Type}}$	WF_TYPE_PAIR
1115	$\frac{\Theta \vdash t \text{ Type} \quad \Theta \vdash t' \text{ Type}}{\Theta \vdash t \multimap t' \text{ Type}}$	WF_TYPE_LOLLY
1116		

1117 **E** NumLin Grammar

	m	::=		matrix expressions
			M	matrix variables
			$m + m'$	matrix addition
			$m m'$	matrix multiplication
1118			(m)	S
	f	::=		fractional permission
			fc	variable
			1	whole permission
			$\frac{1}{2}f$	

23:38 NumLin: Linear Types for Linear Algebra

1119	t	::=		linear type
			unit	unit
			bool	boolean (true/false)
			int	63-bit integers
			elt	array element
			f arr	arrays
			f mat	matrices
			!t	multiple-use type
			$\forall fc.t$	frac. perm. generalisation
			$t \otimes t'$	pair
			$t \multimap t'$	linear function
			(t)	S parentheses
	p	::=		primitive
			not	boolean negation
			$(+)$	integer addition
			$(-)$	integer subtraction
			$(*)$	integer multiplication
			$(/)$	integer division
			$(=)$	integer equality
			$(<)$	integer less-than
			$(+.)$	element addition
			$(-.)$	element subtraction
			$(*.)$	element multiplication
			$(/.)$	element division
			$(=.)$	element equality
			$(<.)$	element less-than
			set	array index assignment
			get	array indexing
			share	share array
			unshare	unshare array
			free	free array
			array	Owl: make array
			copy	Owl: copy array
			sin	Owl: map sine over array
			hypot	Owl: $x_i := \sqrt{x_i^2 + y_i^2}$
			asum	BLAS: $\sum_i x_i $
			axpy	BLAS: $x := \alpha x + y$
			dot	BLAS: $x \cdot y$
			rotmg	BLAS: see its docs
			scal	BLAS: $x := \alpha x$
			amax	BLAS: $\operatorname{argmax} i : x_i$
			setM	matrix index assignment
			getM	matrix indexing

		shareM		share matrix
		unshareM		unshare matrix
		freeM		free matrix
		matrix		Owl: make matrix
		copyM		Owl: copy matrix
		copyM_to		Owl: copy matrix onto another
		sizeM		dimension of matrix
		trnsp		transpose matrix
		gemm		BLAS: $C := \alpha A^{T?} B^{T?} + \beta C$
		symm		BLAS: $C := \alpha AB + \beta C$
		posv		BLAS: Cholesky decomp. and solve
		potrs		BLAS: solve with given Cholesky
		syrk		BLAS: $C := \alpha A^{T?} A^{T?} + \beta C$
<i>v</i>	::=			values
		<i>p</i>		primitives
		<i>x</i>		variable
		()		unit introduction
		true		true
		false		false
		<i>k</i>		integer
		<i>l</i>		heap location
		<i>el</i>		array element
		Many <i>v</i>		!-introduction
		fun <i>fc</i> $\rightarrow v$		frac. perm. abstraction
		(<i>v</i> , <i>v'</i>)		pair introduction
		fun <i>x</i> : <i>t</i> $\rightarrow e$	bind <i>x</i> in <i>e</i>	abstraction
		fix (<i>g</i> , <i>x</i> : <i>t</i> , <i>e</i> : <i>t'</i>)	bind <i>g</i> $\cup x$ in <i>e</i>	fixpoint
		(<i>v</i>)	S	parentheses
<i>e</i>	::=			expression
		<i>p</i>		primitives
		<i>x</i>		variable
		let <i>x</i> = <i>e</i> in <i>e'</i>	bind <i>x</i> in <i>e'</i>	let binding
		()		unit introduction
		let () = <i>e</i> in <i>e'</i>		unit elimination
		true		true
		false		false
		if <i>e</i> then <i>e</i> ₁ else <i>e</i> ₂		if
		<i>k</i>		integer
		<i>l</i>		heap location
		<i>el</i>		array element
		Many <i>e</i>		!-introduction
		let Many <i>x</i> = <i>e</i> in <i>e'</i>		!-elimination
		fun <i>fc</i> $\rightarrow e$		frac. perm. abstraction

23:40 NumLin: Linear Types for Linear Algebra

		$e[f]$		frac. perm. specialisation
		(e, e')		pair introduction
		$\mathbf{let} (a, b) = e \mathbf{in} e'$	bind $a \cup b$ in e'	pair elimination
		$\mathbf{fun} x : t \rightarrow e$	bind x in e	abstraction
		$e e'$		application
		$\mathbf{fix} (g, x : t, e : t')$	bind $g \cup x$ in e	fixpoint
		(e)	S	parentheses
C	$::=$			evaluation contexts
		$\mathbf{let} x = [-] \mathbf{in} e$	bind x in e	let binding
		$\mathbf{let} () = [-] \mathbf{in} e$		unit elimination
		$\mathbf{if} [-] \mathbf{then} e_1 \mathbf{else} e_2$		if
		$\mathbf{Many} [-]$!-introduction
		$\mathbf{let Many} x = [-] \mathbf{in} e$!-elimination
		$\mathbf{fun} fc \rightarrow [-]$		frac. perm. abstraction
		$[-][f]$		frac. perm. specialisation
		$([-], e)$		pair introduction
		$(v, [-])$		pair introduction
		$\mathbf{let} (a, b) = [-] \mathbf{in} e$	bind $a \cup b$ in e	pair elimination
		$[-]e$		application
1121		$v[-]$		application
Θ	$::=$			fractional permission environment
		\cdot		
		Θ, fc		
Γ	$::=$			linear types environment
		\cdot		
		$\Gamma, x : t$		
		Γ, Γ'		
Δ	$::=$			intuitionistic types environment
		\cdot		
		$\Delta, x : t$		
σ	$::=$			heap (multiset of triples)
		$\{\}$		empty heap
		$\sigma + \{l \mapsto_f m_{k_1, k_2}\}$		location l points to matrix m
$Config$	$::=$			result of small step
		$\langle \sigma, e \rangle$		heap and expression
		\mathbf{err}		error

1122

F Desugaring NumLin

$$\begin{aligned}
x[e] &\Rightarrow \text{get } _ x (e) && \text{(similarly for matrices)} \\
x[e_1] := e_2 &\Rightarrow \text{set } x (e_1) (e_2) && \text{(similarly for matrices)}
\end{aligned}$$

$$\begin{aligned}
pat &::= () \mid x \mid !x \mid \text{Many } pat \mid (pat, pat) \\
\text{let } !x = e_1 \text{ in } e_2 &\Rightarrow \text{let } \text{Many } x = e_1 \text{ in} \\
&\quad \text{let } \text{Many } x = \text{Many } (\text{Many } x) \text{ in } e_2 \\
\text{let } \text{Many } \langle pat_x \rangle = e_1 \text{ in } e_2 &\Rightarrow \text{let } \text{Many } x = x \text{ in} \\
&\quad \text{let } \langle pat_x \rangle = x \text{ in } e_2 \\
\text{let } (\langle pat_a \rangle, \langle pat_b \rangle) = e_1 \text{ in } e_2 &\Rightarrow \text{let } (a, b) = a_b \text{ in let } \langle pat_a \rangle = a \text{ in} \\
&\quad \text{let } \langle pat_b \rangle = b \text{ in } e_2 \\
\text{fun } (\langle pat_x \rangle : t) \rightarrow e &\Rightarrow \text{fun } (x : t) \rightarrow \text{let } \langle pat_x \rangle = x \text{ in } e
\end{aligned}$$

1123

$$\begin{aligned}
arg &::= \langle pat \rangle : t \mid 'x \text{ (fractional permission variable)} \\
\text{fun } \langle arg_{1..n} \rangle \rightarrow e &\Rightarrow \text{fun } \langle arg_1 \rangle \rightarrow .. \text{fun } \langle arg_n \rangle \rightarrow e \\
\text{let } f \langle arg_{1..n} \rangle = e_1 \text{ in } e_2 &\Rightarrow \text{let } f = \text{fun } \langle arg_{1..n} \rangle \rightarrow e_1 \text{ in } e_2 \\
\text{let } !f \langle arg_{1..n} \rangle = e_1 \text{ in } e_2 &\Rightarrow \text{let } \text{Many } f = \text{Many } (\text{fun } \langle arg_{1..n} \rangle \rightarrow e_1) \text{ in } e_2 \\
\text{fixpoint} &\equiv \text{fix } (f, x : t, \text{fun } \langle arg_{1..n} \rangle \rightarrow e_1 : t') \\
\text{let rec } f (x : t) \langle arg_{1..n} \rangle : t' = e_1 \text{ in } e_2 &\Rightarrow \text{let } f = \text{fixpoint in } e_2 \\
\text{let rec } !f (x : t) \langle arg_{1..n} \rangle : t' = e_1 \text{ in } e_2 &\Rightarrow \text{let } \text{Many } f = \text{Many } \text{fixpoint in } e_2
\end{aligned}$$

G Primitives

```

module Arr = Owl.Dense.Ndarray.D
type z = Z
type 'a s = Succ
type 'a arr = A of Arr.arr [@@unboxed]
type 'a mat = M of Arr.arr [@@unboxed]
type 'a bang = Many of 'a [@@unboxed]
module Prim :
sig
  val extract : 'a bang -> 'a
  (** Boolean *)
  val not_ : bool bang -> bool bang
  (** Arithmetic, many omitted for brevity *)
  val addI : int bang -> int bang -> int bang
  val eqI : int bang -> int bang -> bool bang
  (** Arrays *)
  val set : z arr -> int bang -> float bang -> z arr
  val get : 'a arr -> int bang -> 'a arr * float bang
  val share : 'a arr -> 'a s arr * 'a s arr
  val unshare : 'a s arr -> 'a s arr -> 'a arr
  val free : z arr -> unit
  (** Owl *)
  val array : int bang -> z arr
  val copy : 'a arr -> 'a arr * z arr
  val sin : z arr -> z arr
  val hypot : z arr -> 'a arr -> 'a arr * z arr
  (** Level 1 BLAS *)
  val asum : 'a arr -> 'a arr * float bang
  val axpy : float bang -> 'a arr -> z arr -> 'a arr * z arr
  val dot : 'a arr -> 'b arr -> ('a arr * 'b arr) * float bang
  val scal : float bang -> z arr -> z arr
  val amax : 'a arr -> 'a arr * int bang
  (* Matrix, some omitted for brevity *)
  val matrix : int bang -> int bang -> z mat
  val eye : int bang -> z mat
  val copy_mat : 'a mat -> 'a mat * z mat
  val copy_mat_to : 'a mat -> z mat -> 'a mat * z mat
  val size_mat : 'a mat -> 'a mat * (int bang * int bang)
  val transpose : 'a mat -> 'a mat * z mat
  (* Level 3 BLAS/LAPACK *)
  val gemm : float bang -> ('a mat * bool bang) -> ('b mat * bool bang) ->
    float bang -> z mat -> ('a mat * 'b mat) * z mat
  val symm : bool bang -> float bang -> 'a mat -> 'b mat -> float bang ->
    z mat -> ('a mat * 'b mat) * z mat
  val gesv : z mat -> z mat -> z mat * z mat
  val posv : z mat -> z mat -> z mat * z mat
  val potrs : 'a mat -> z mat -> 'a mat * z mat
  val syrks : bool bang -> float bang -> 'a mat -> float bang -> z mat ->
    'a mat * z mat
end

```

1126

H Kalman Filters from NumLin and C

```

let kalman sigma h mu r_1 data_1 =
  let h, _p_k_n_p_ = Prim.size_mat h in
  let k, n = _p_k_n_p_ in
  let sigma_hT = Prim.matrix n k in
  let (sigma, h), sigma_hT =
    Prim.gemm (Many 1.) (sigma, Many false) (h, Many true) (Many 0.) sigma_hT in
  let (h, sigma_hT), r_2 =
    Prim.gemm (Many 1.) (h, Many false) (sigma_hT, Many false) (Many 1.) r_1 in
  let k_by_k, x = Prim.posv_flip r_2 sigma_hT in
  let (h, mu), data_2 =
    Prim.gemm (Many 1.) (h, Many false) (mu, Many false) (Many (-1.)) data_1 in
  let (x, data_2), new_mu =
    Prim.gemm (Many 1.) (x, Many false) (data_2, Many false) (Many 1.) mu in
  let x_h = Prim.matrix n n in
  let (x, h), x_h =
    Prim.gemm (Many 1.) (x, Many false) (h, Many false) (Many 0.) x_h in
  let () = Prim.free_mat x in
  let sigma, sigma2 = Prim.copy_mat sigma in
  let (sigma, x_h), new_sigma =
    Prim.symm (Many true) (Many (-1.)) sigma x_h (Many 1.) sigma2 in
  let () = Prim.free_mat x_h in
  ((sigma, h), (new_sigma, (new_mu, (k_by_k, data_2))))

```

■ **Figure 16** OCaml code for a Kalman filter, generated (at *compile time*) from the code in Figure 8, passed through `ocamlformat` for presentation.

```

static void kalman( const int n,          const int k,
                  const double *sigma, /* n,n */ const double *h, /* k,n */
                  const double *mu, /* n,1 */ double *r, /* k,k */
                  double *data, /* k,1 */ double **ret_sigma /* n,n */ ) {
  double* n_by_k = (double *) malloc(n * k * sizeof(double));
  cblas_dgemm(RowMajor, NoTrans, Trans, n, k, n, 1., sigma, n, h, n, 0., n_by_k, k);
  cblas_dgemm(RowMajor, NoTrans, NoTrans, k, k, n, 1., h, n, n_by_k, k, 1., r, k);
  LAPACKE_dposv(LAPACK_COL_MAJOR, 'U', k, n, r, k, n_by_k, k);
  cblas_dgemm(RowMajor, NoTrans, NoTrans, k, 1, n, 1., h, n, mu, 1, -1., data, 1);
  cblas_dgemm(RowMajor, NoTrans, NoTrans, n, 1, k, 1., n_by_k, k, data, 1, 1., mu, 1);
  double* n_by_n = (double *) malloc(n * n * sizeof(double));
  cblas_dgemm(RowMajor, NoTrans, NoTrans, n, n, k, 1., n_by_k, k, h, n, 0., n_by_n, n);
  free(n_by_k);
  double* n_by_n2 = (double *) malloc(n * n * sizeof(double));
  cblas_dcopy(n*n, sigma, 1, n_by_n2, 1);
  cblas_dsymm(RowMajor, Right, Upper, n, n, -1., sigma, n, n_by_n, n, 1., n_by_n2, n);
  free(n_by_n);
  *ret_sigma = n_by_n2; }

```

■ **Figure 17** CBLAS/LAPACKE implementation of a Kalman filter. I used C instead of Fortran because it is what Owl uses under the hood and OCaml FFI support for C is better and easier to use than that for Fortran. A distinct ‘`measure_kalman`’ function that sandwiches a call to this function with `getrusage` is omitted for brevity.